

Do you trust profilers? I once did too

Johannes Bechberger



Nightmares of a profiler developer



Johannes Bechberger's gruesome tales on Java profiling APIs

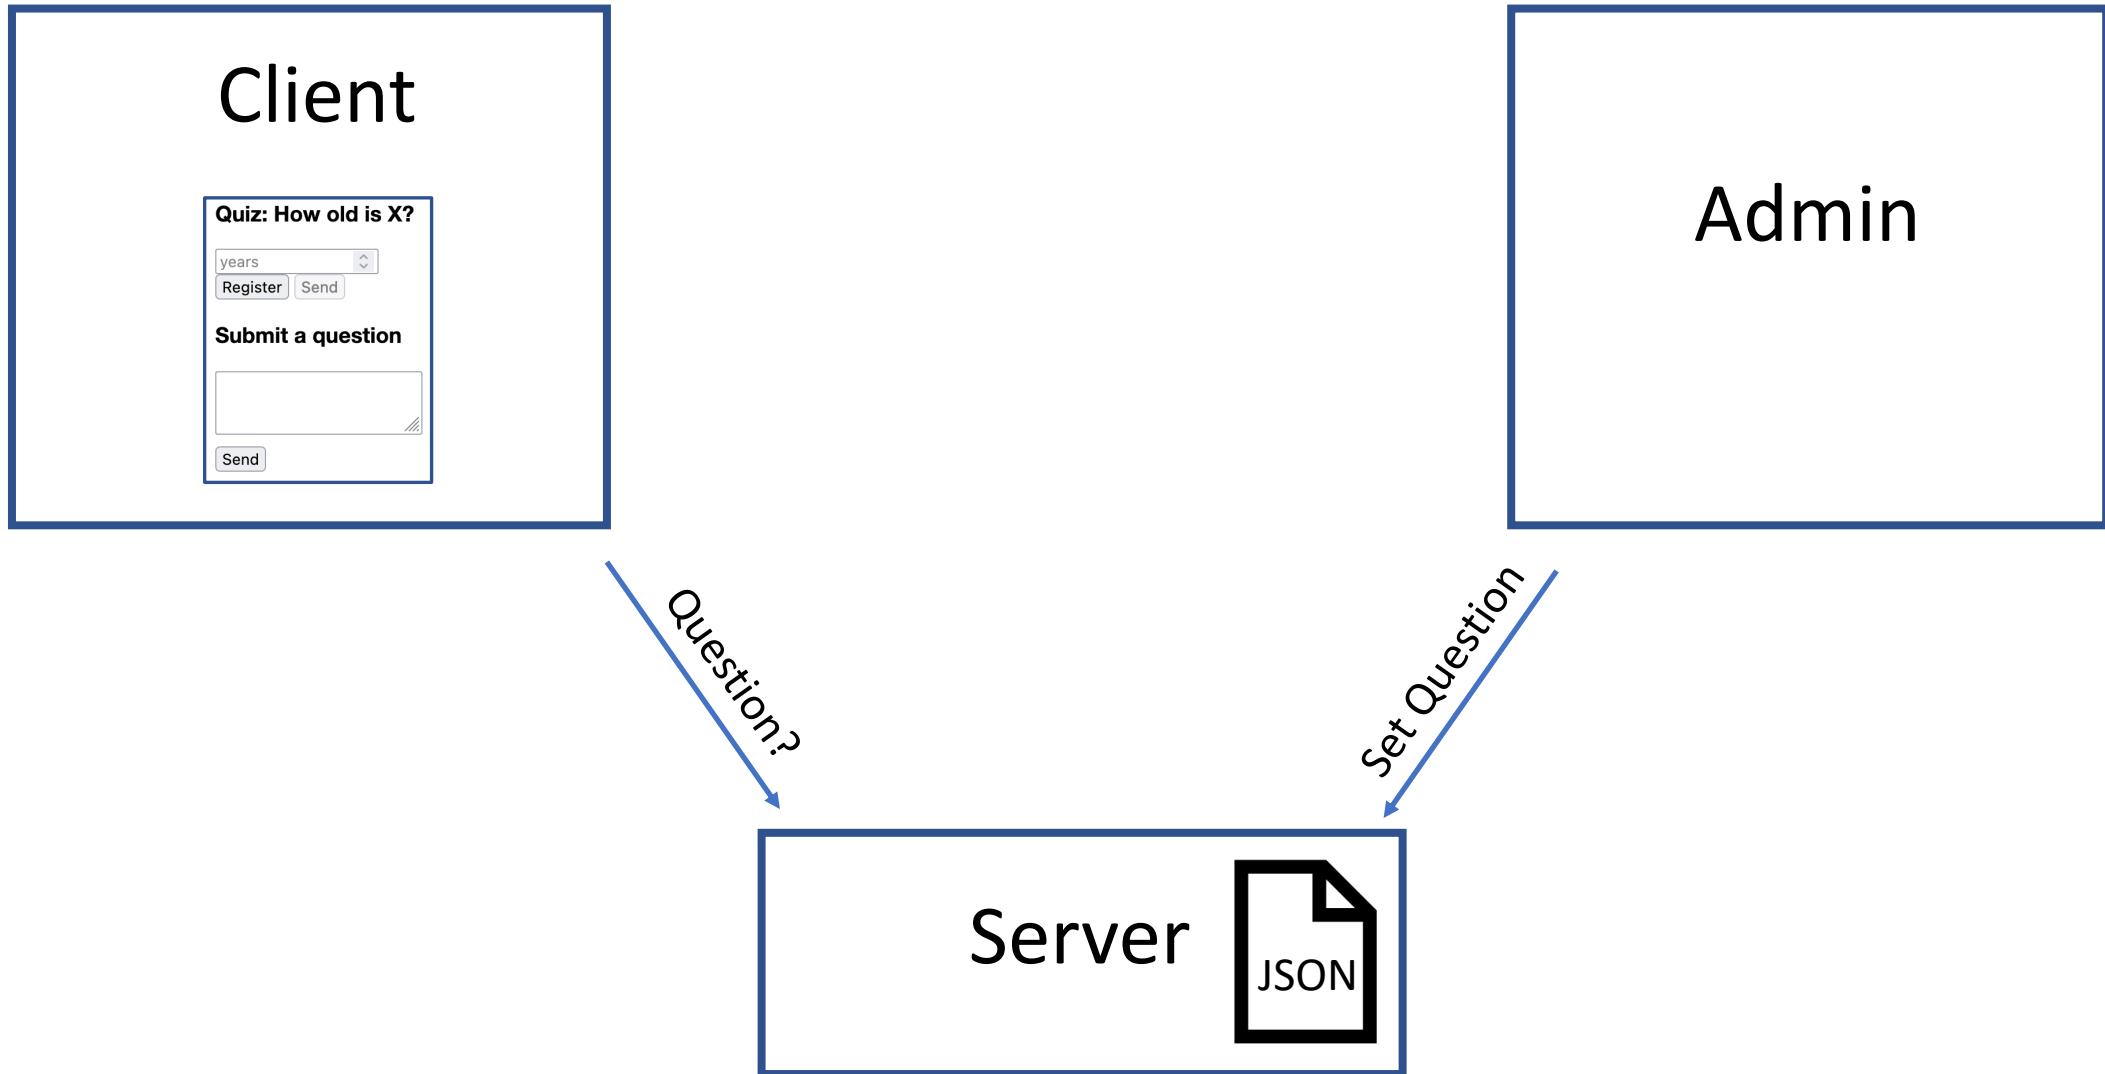
Quiz: How old is X?

years

Register Send

Submit a question

Send



Premature Optimization

“

premature
optimization is the root of all evil.



Source: Jacob Appelbaum, Wikipedia

— Donald Knuth

So the event came...

... and the server was overwhelmed

Premature Optimization

“ We *should* forget about small efficiencies, say about 97% of the time: premature optimization is the root of all evil.

Yet we should not pass up our opportunities in that critical 3%.

A good programmer will not be lulled into complacency by such reasoning, he will be wise to look carefully at the critical code; **but only after that code has been identified.**

— Donald Knuth



Source: Jacob Appelbaum, Wikipedia

Profilers to the rescue

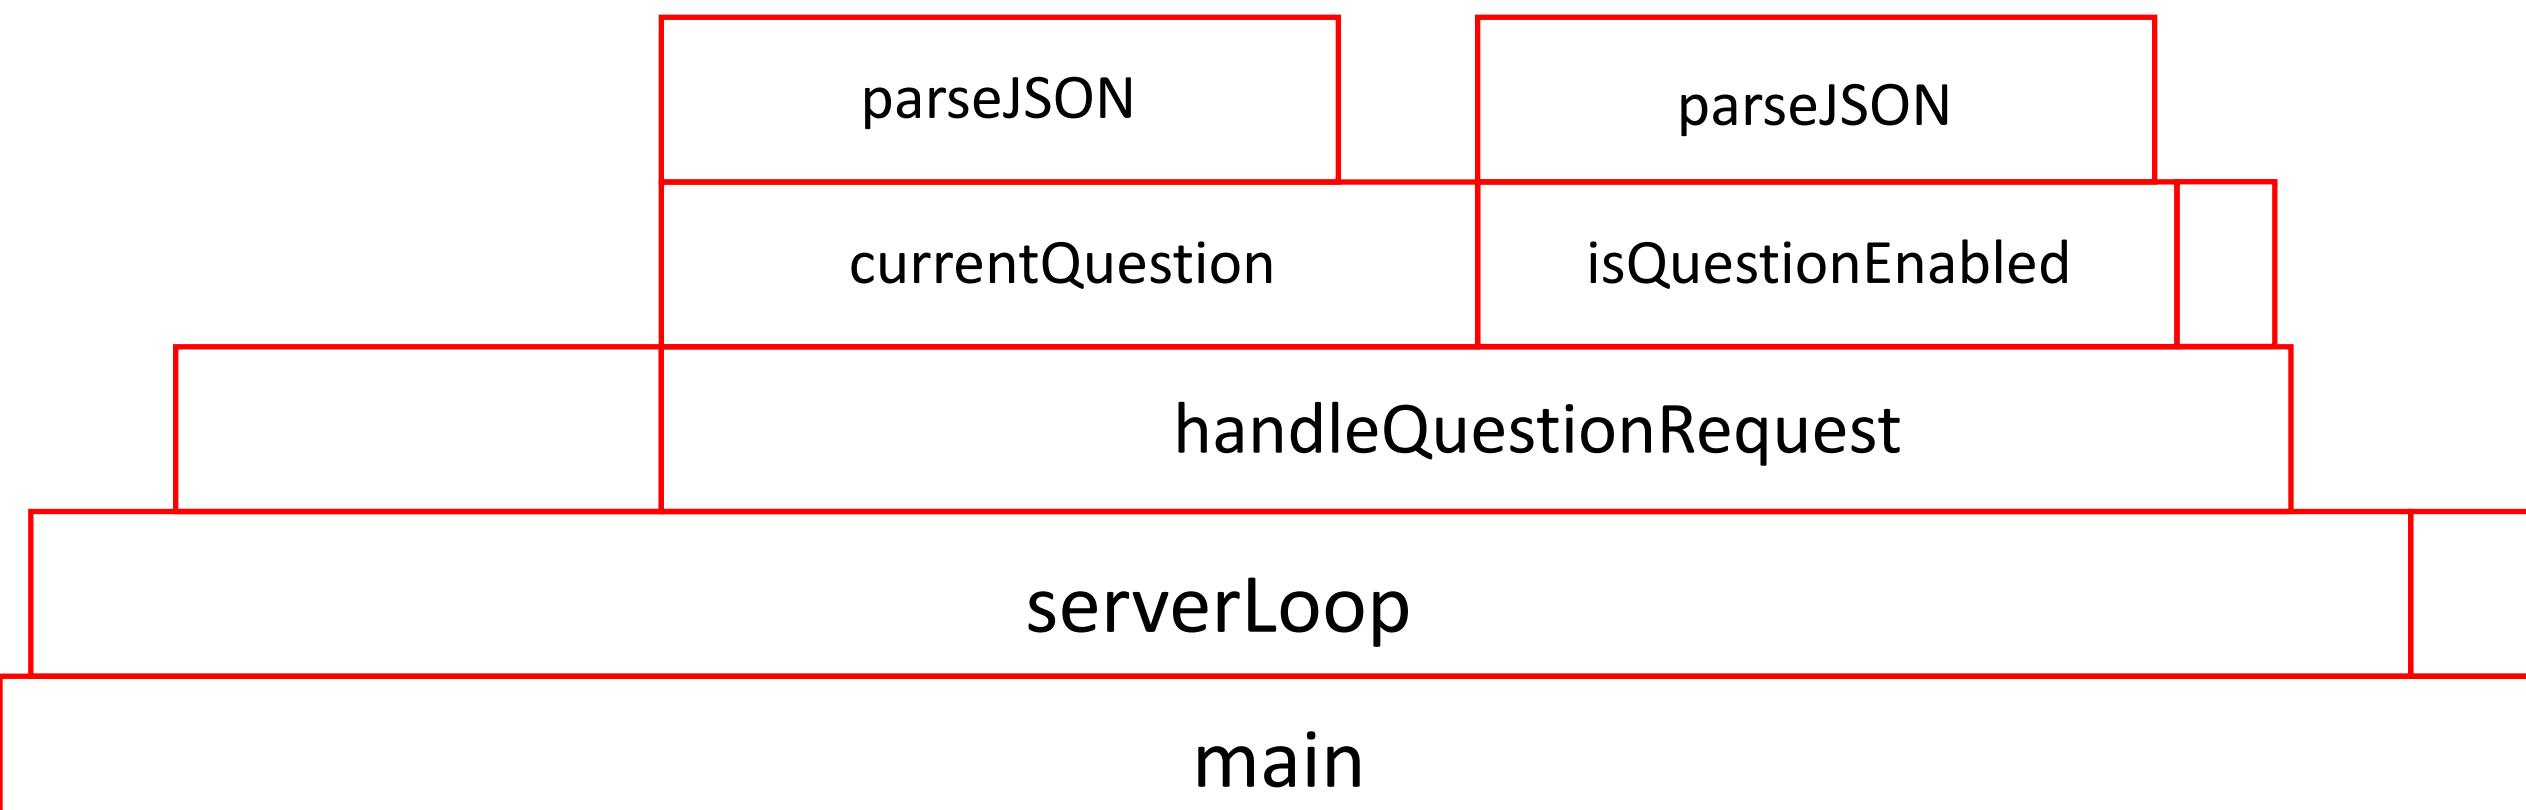




Source: Abgedreht Karlsruhe

```
fun main() {  
    ...  
    serverLoop()  
}  
  
    fun serverLoop() {  
        while (true) {  
            req = ...  
            if (req.isQuestionRequest)  
                handleQuestionRequest(req)  
            ...  
        }  
    }  
}
```

```
    fun handleQuestionRequest(req) {  
        if (isQuestionEnabled()) {  
            emit(currentQuestion().json)  
        } else {  
            emit("{}")  
        }  
    }  
}
```





Mario Fusco 🇪🇺🇺🇦 @mariofusco@jvm.social

@mariofusco

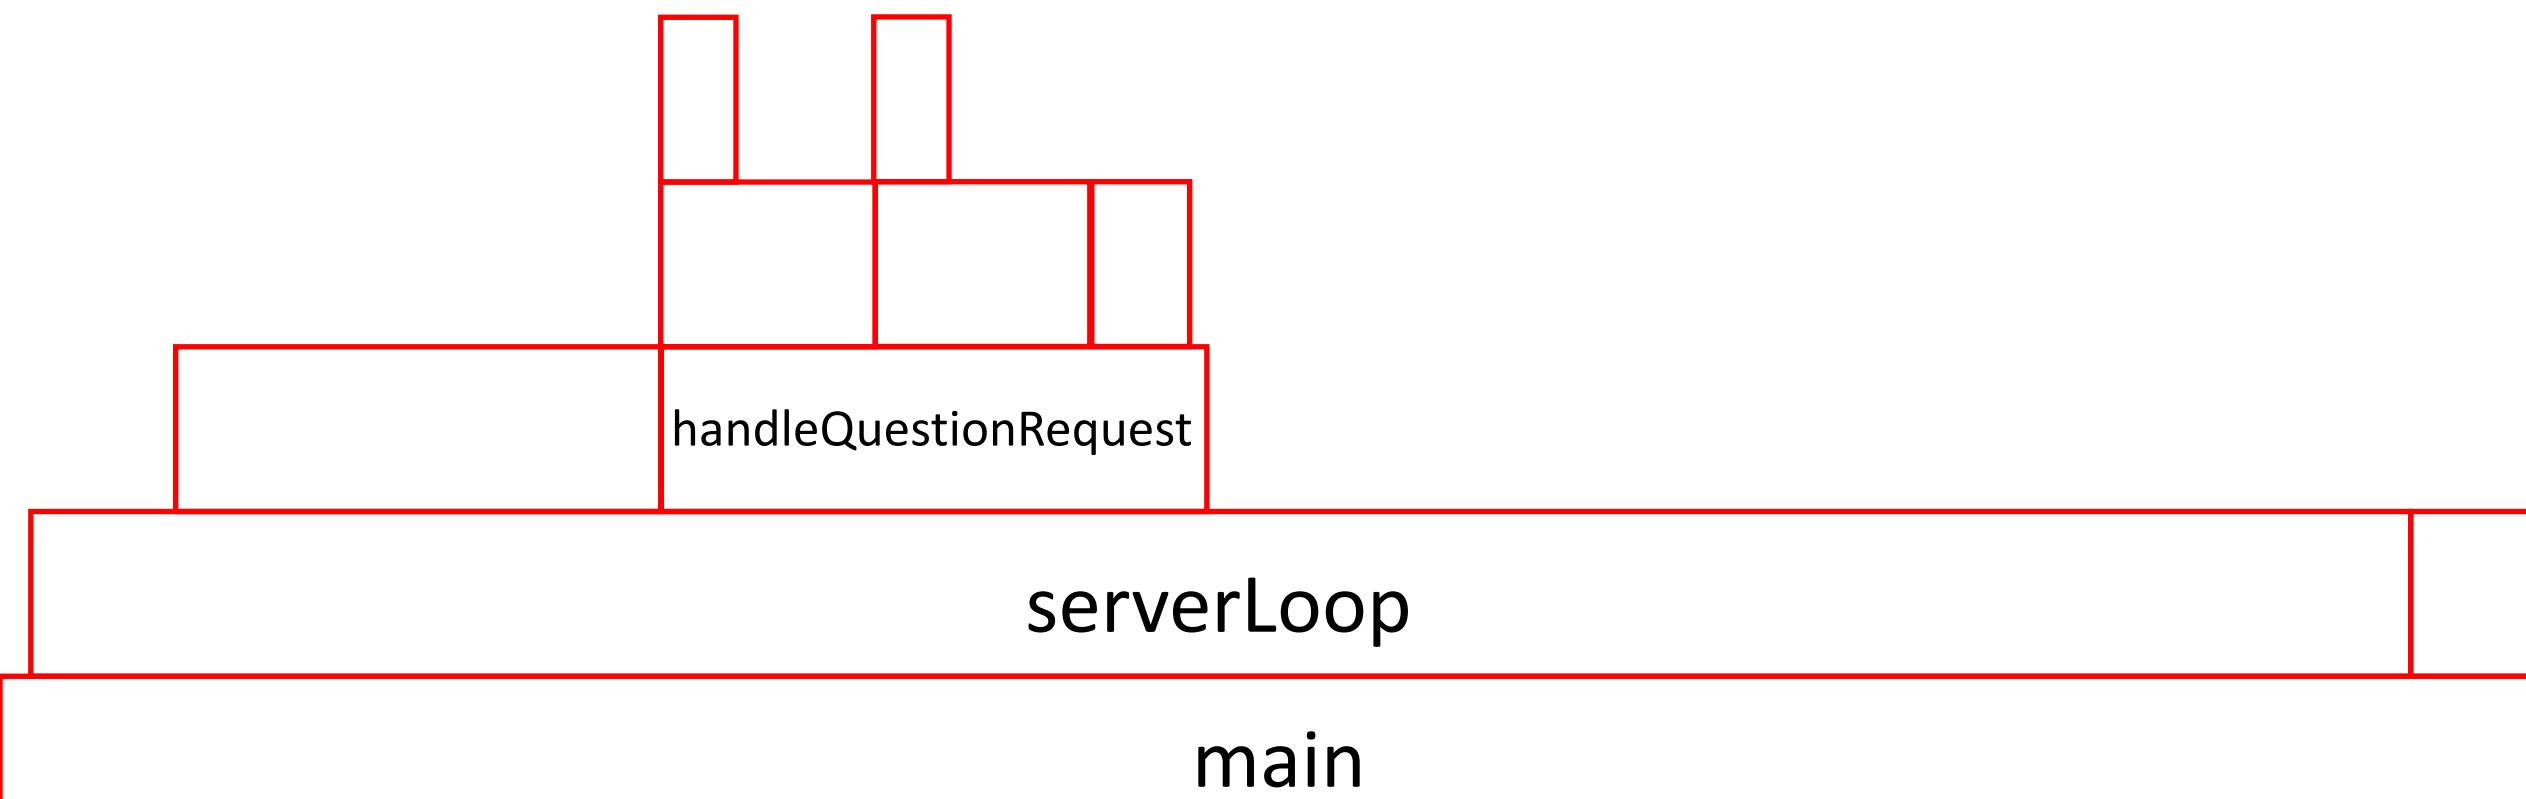
...

I love flamegraphs. When you do something stupid, it punches in your face and it's impossible not see it.

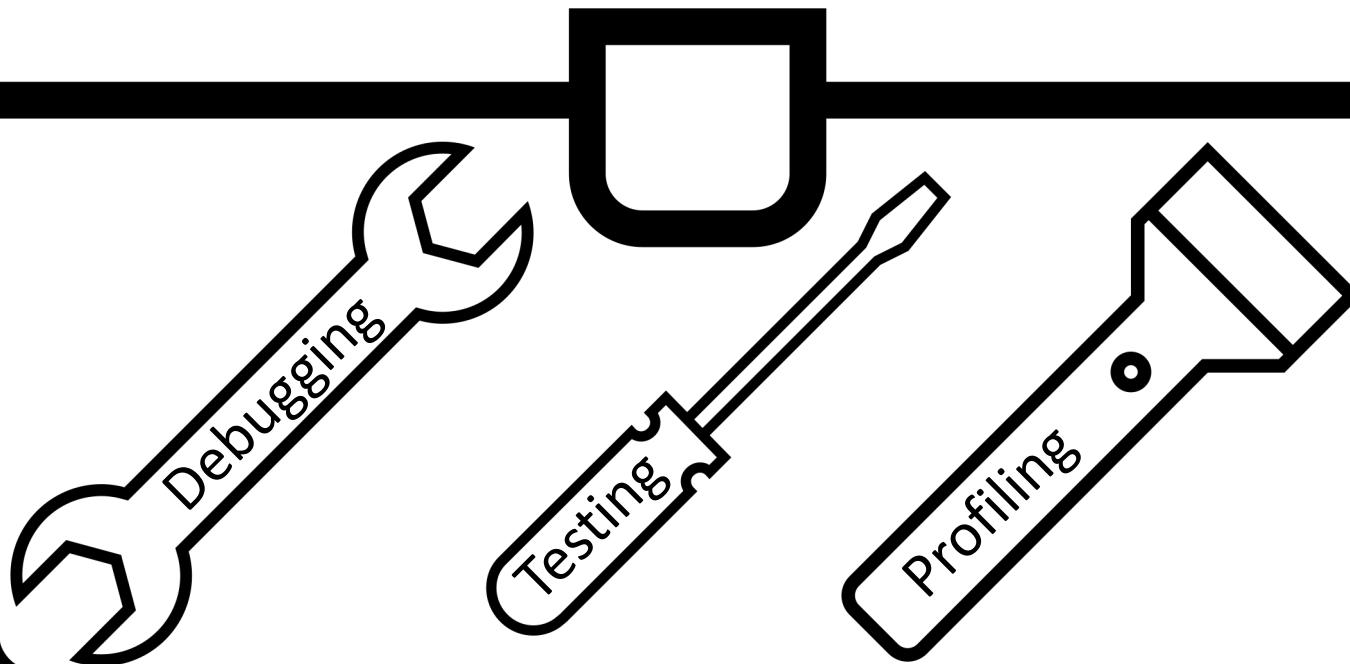


```
fun main() {  
    ...  
    serverLoop()  
}  
  
    fun serverLoop() {  
        while (true) {  
            req = ...  
            if (req.isQuestionRequest)  
                handleQuestionRequest(req)  
            ...  
        }  
    }  
}
```

```
    fun handleQuestionRequest(req) {  
        if (isQuestionEnabled()) {  
            emit(currentQuestion().json)  
        } else {  
            emit({})  
        }  
    }  
}
```

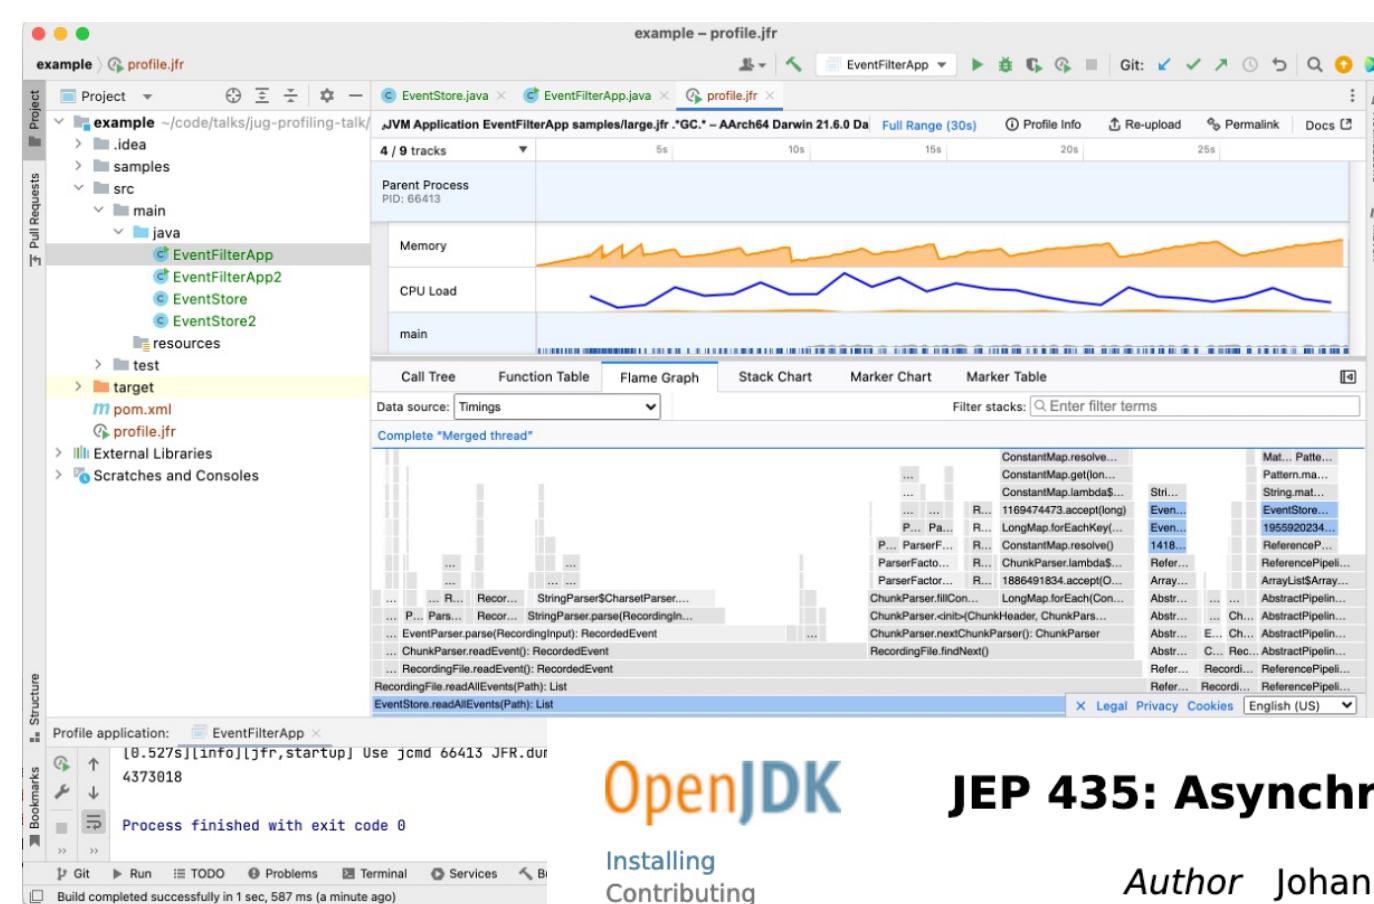


Toolbox



jar profiler





[Installing](#)

[Contributing](#)

[Sponsoring](#)

[Developers' Guide](#)

[Vulnerabilities](#)

[JDK GA/EA Builds](#)

[Mailing lists](#)

[Wiki · IRC](#)

[Bylaws · Census](#)

[Legal](#)

[JEP Process](#)

[Source code](#)

[Mercurial](#)

[GitHub](#)

[Tools](#)

[Git](#)

[jtrex harness](#)

JEP 435: Asynchronous Stack Trace VM API

Author Johannes Bechberger

Owner Christoph Langer

Type Feature

Scope JDK

Status Candidate

Component hotspot / svc

Discussion serviceability dash dev at openjdk dot org

Effort S

Duration S

Reviewed by Andrei Pangin, Christoph Langer, Jaroslav Bachorík

Created 2022/04/04 11:02

Updated 2023/01/19 13:08

async-profiler

This project is a low overhead sampling profiler for Java that does not suffer from the memory overhead of HotSpot. It uses features HotSpot-specific APIs to collect stack traces and to track memory usage. It works with OpenJDK, Oracle JDK and other Java runtimes based on the HotSpot JVM.

async-profiler can trace the following kinds of events:

- CPU cycles
- Hardware and Software performance counters like cache misses, branch switches, etc.

Who of you used a
profiler?

THE NEW
HACKER'S
DICTIONARY

second edition



foreword and cartoons by Guy L. Steele Jr.

compiled by

ERIC S. RAYMOND

What is profiling?

“

A report on the amounts of time spent in each routine of a program, used to find and tune away the hot spots in it.

— The Jargon File

Different Profilers

Instrumenting Profilers

Inserting instructions into the code automatically

```
fun serverLoop() {  
    logEntry("serverLoop")  
    while (true) {  
        req = ...  
        if (req.isQuestionRequest)  
            handleQuestionRequest(req)  
        ...  
    }  
    logExit("serverLoop")  
}
```

But JITs?

Sampling Profilers

```
fun serverLoop() {  
    while (true) {  
        req = ...  
        if (req.isQuestionRequest)  
            handleQuestionRequest(req)  
        ...  
    }  
}
```

Profilers aren't
rocket science

So let's write our own*

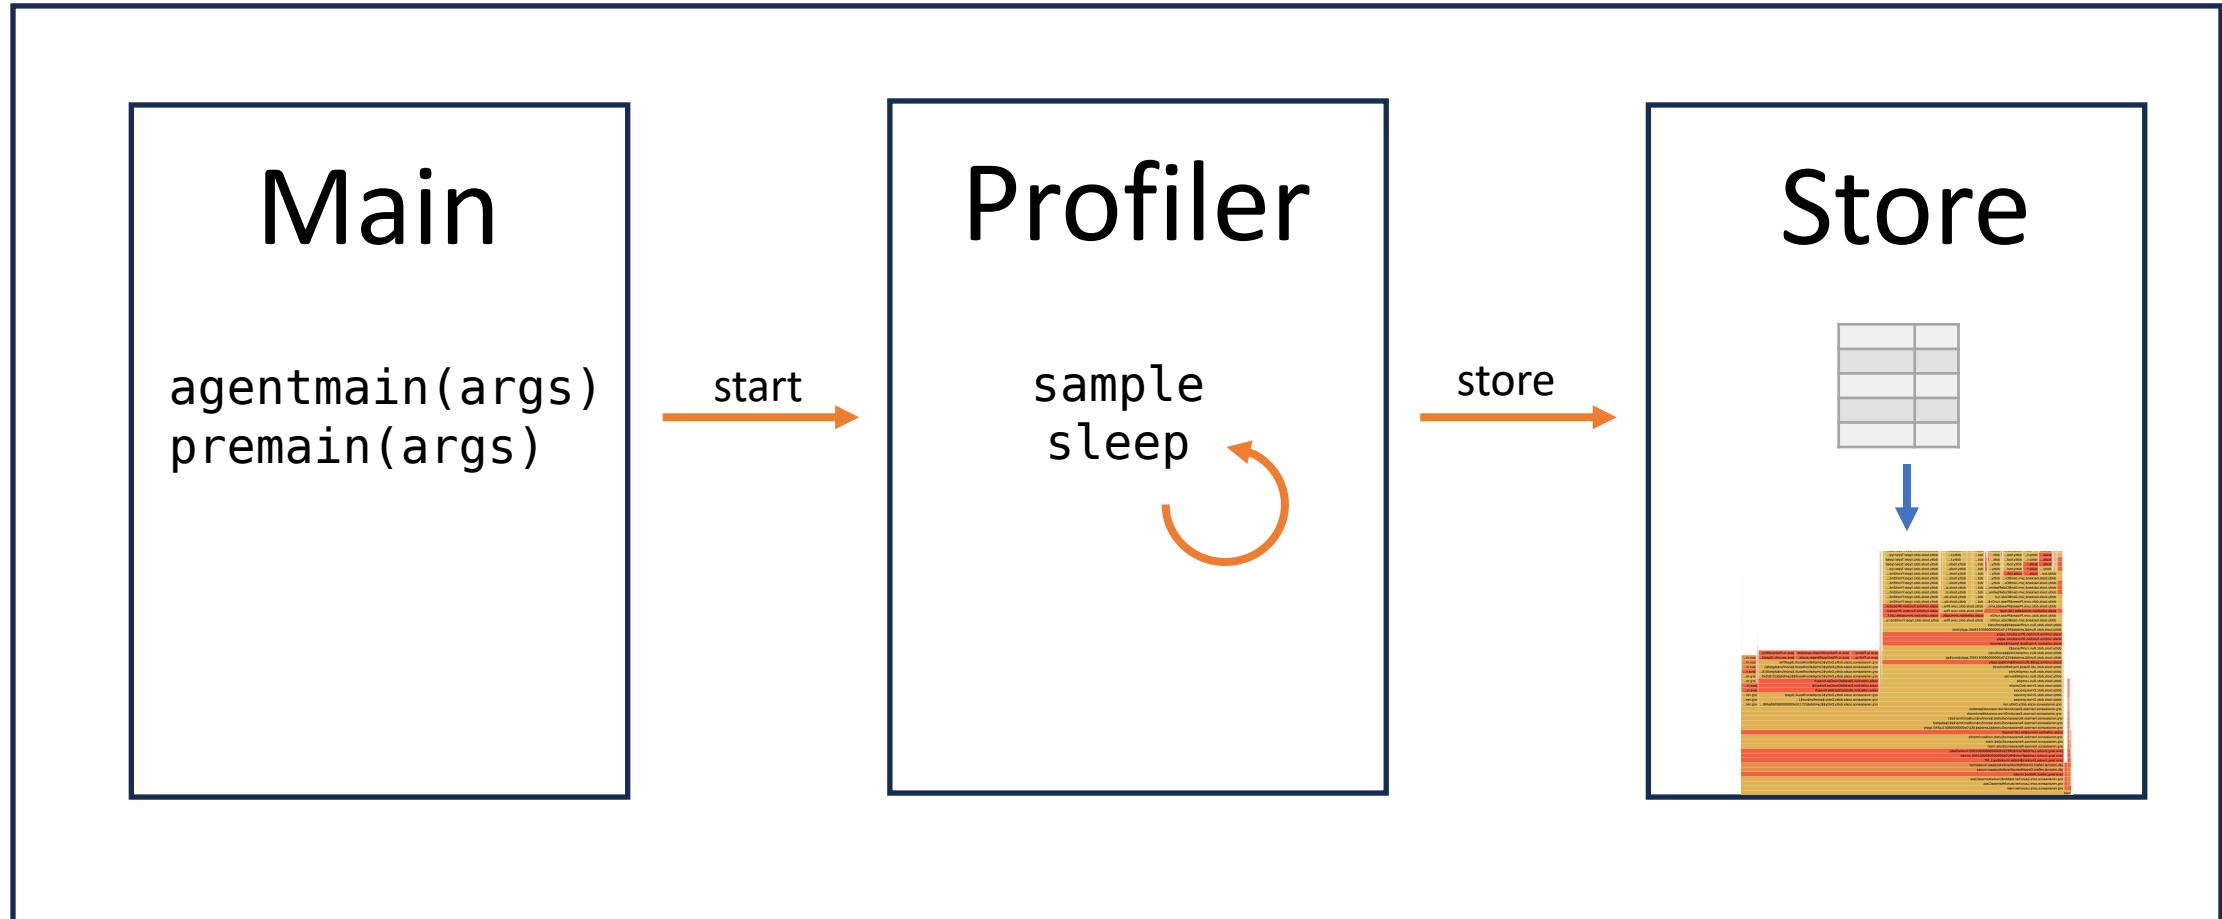
*ignoring safepoint bias

Writing a Profiler in 240 Lines of Pure Java

Posted on March 27, 2023

A few months back, I started writing a profiler from scratch, and the code since became the base of my profiler validation tools. The only problem with this project: I wanted to write a proper non-safepoint-biased profiler from scratch. This is a noble effort, but it requires lots C/C++/Unix programming which is finicky, and not everyone can read C/C++ code.

<https://mostlynerdless.de>



Main

agentmain(args)
premain(args)

Main Class

```
java -javaagent:./target/tiny_profiler.jar=agentArgs ...
```

Normally

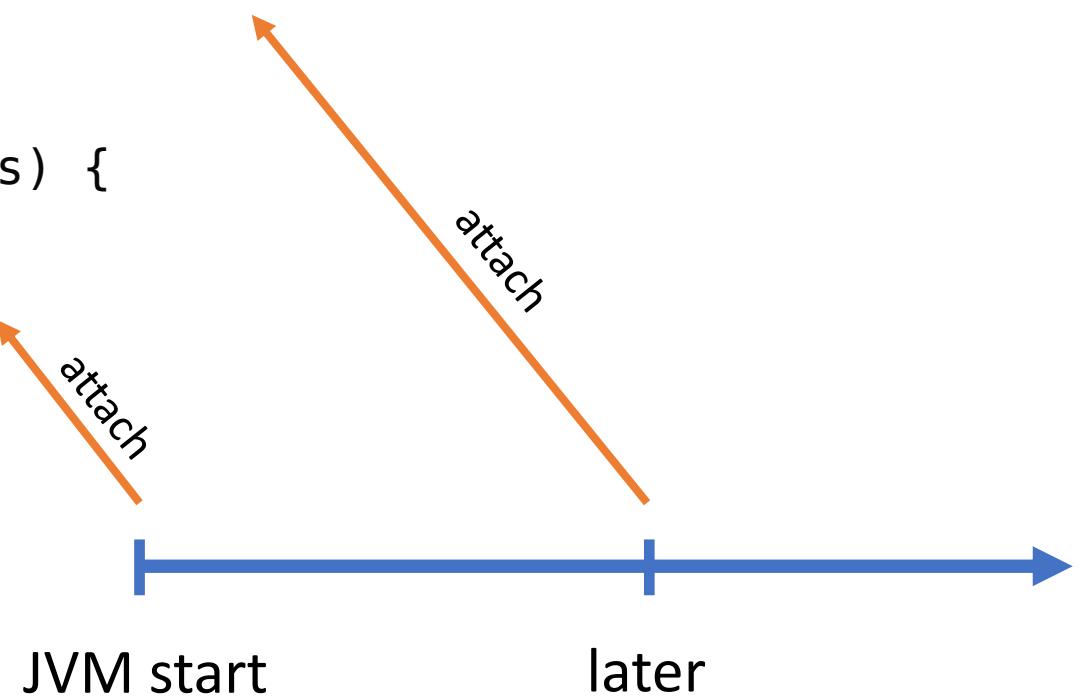
```
public static void main(String[] args) {  
    ...  
}  
  
public static void main(String args[]) {  
    ...  
}  
  
void main() {  
    ...  
}
```

Main Class

```
java -javaagent:./target/tiny_profiler.jar=agentArgs ...
```

```
public static void agentmain(String agentArgs) {  
    premain(agentArgs);  
}
```

```
public static void premain(String agentArgs) {  
    Main main = new Main();  
    main.run(new Options(agentArgs));  
}
```



Main class

```
private void run(Options options) {  
  
    Thread t = new Thread(  
        Profiler.newInstance(options));  
  
    t.setDaemon(true);  
    t.setName("Profiler");  
  
    t.start();  
}  
  
public class Options {  
  
    Duration interval;  
  
    Optional<Path> flamePath;  
  
    boolean printMethodTable;  
}
```

Profiler

sample
sleep



Profiler

```
public static Profiler newInstance(Options options) {  
    Profiler profiler = new Profiler(options);  
  
    Runtime.getRuntime()  
        .addShutdownHook(  
            new Thread(profiler::onEnd));  
  
    return profiler;  
}
```

Profiler

```
private void sample() {
    Thread.getAllStackTraces()

        .forEach(
            (thread, st) -> {

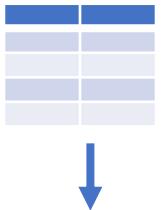
                if (thread.isDaemon()) {
                    return;
                }

                store.addSample(st);
            }
        );
}
```

StackTraceElement

```
class StackTraceElement {  
  
    Class<?> declaringClassObject;  
  
    String classLoaderName;  
    String moduleName;  
    String moduleVersion;  
  
    String declaringClass;  
  
    String methodName;  
    String fileName;  
    int lineNumber;  
  
}
```

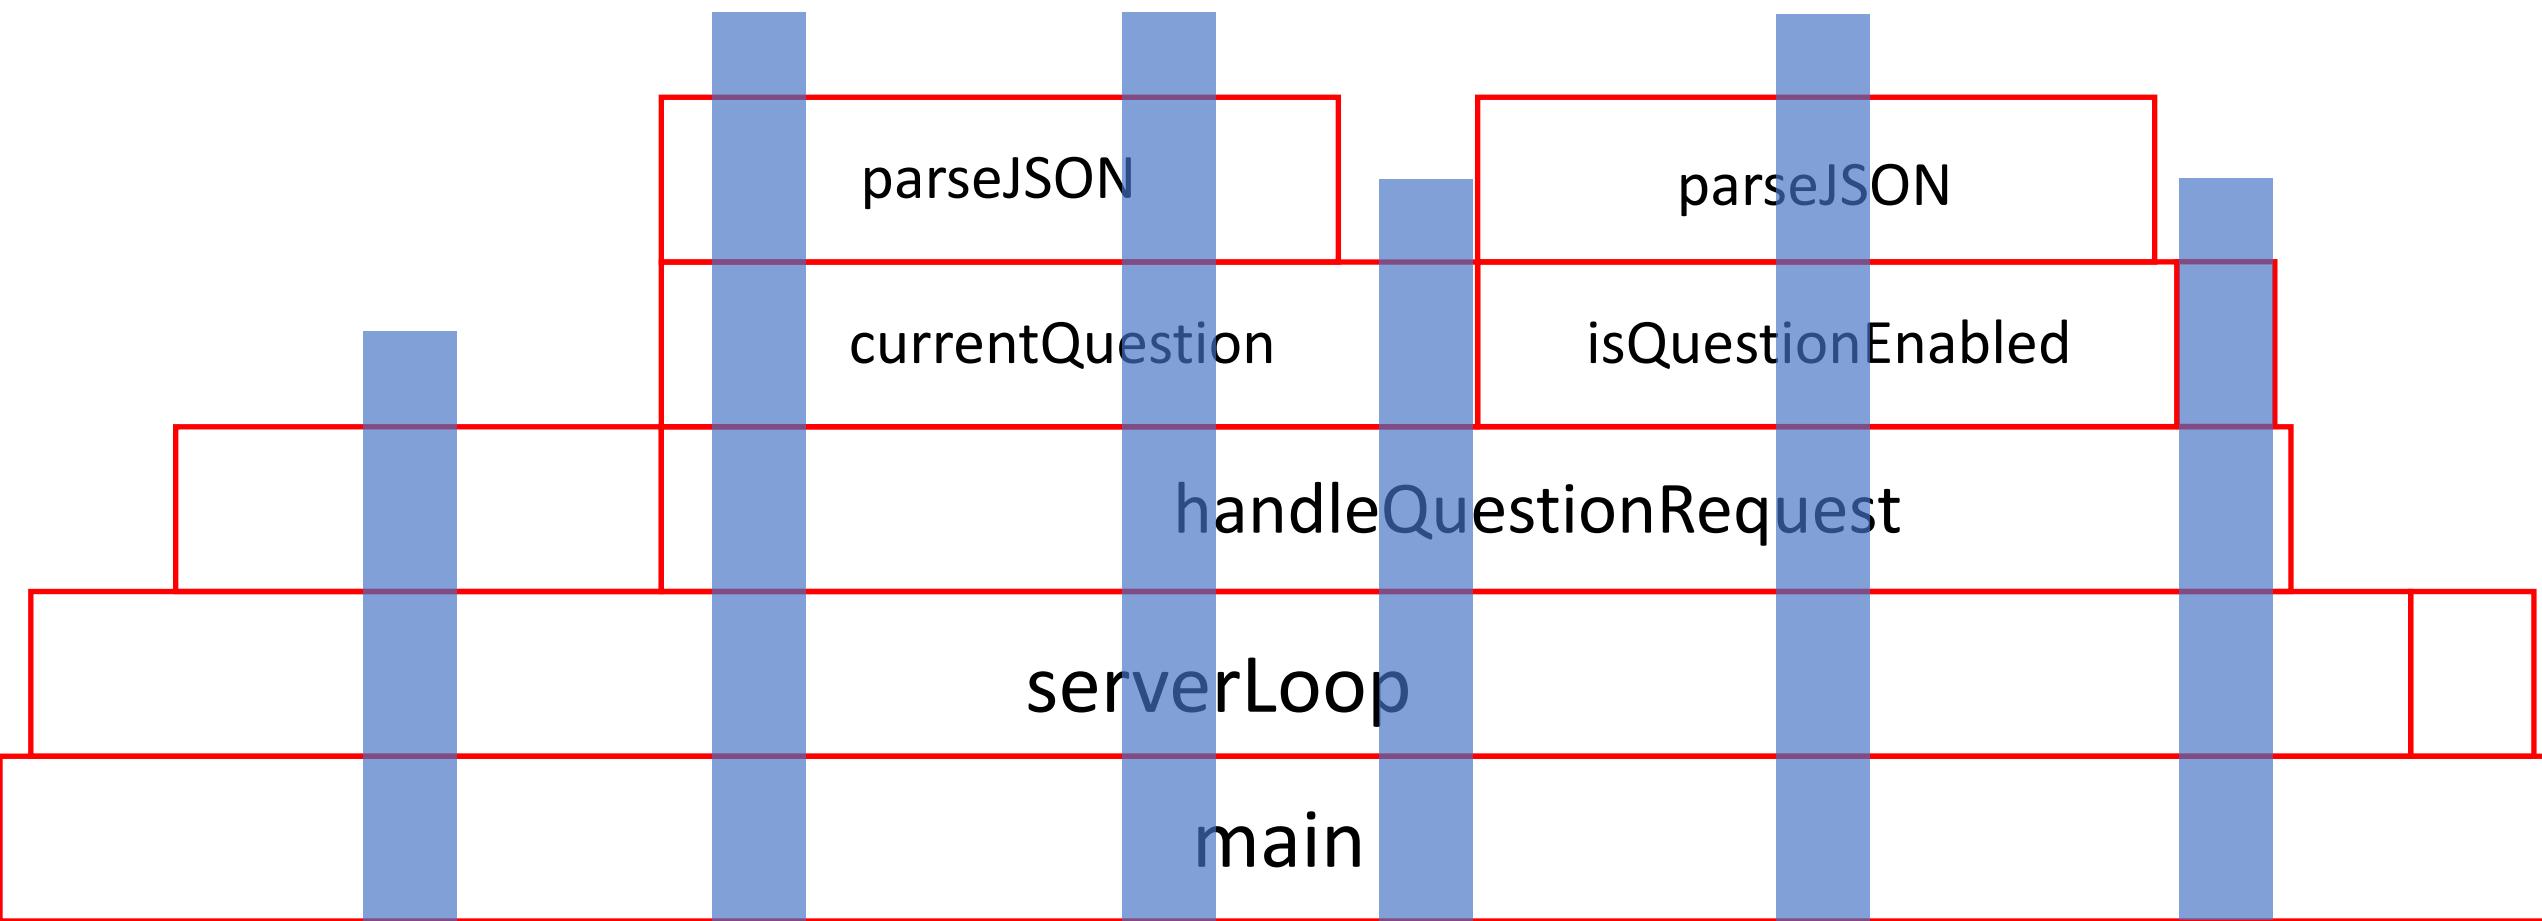
Store



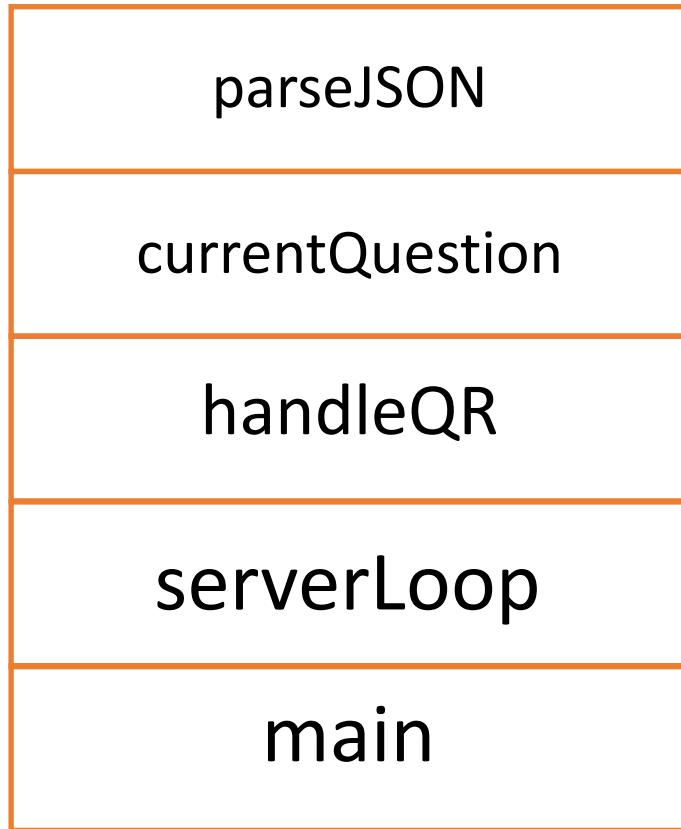
Store

```
class Store {  
  
    Optional<Path> flamePath;  
  
    Map<String, Long> methodOnTopSampleCount;  
    Map<String, Long> methodSampleCount;  
    long totalSampleCount = 0;  
  
    Node rootNode = new Node("root");  
  
    void addSample(StackTraceElement[] st)  
  
    void printMethodTable()  
  
    void storeFlameGraphIfNeeded()  
  
}
```

Flamegraphs would be great

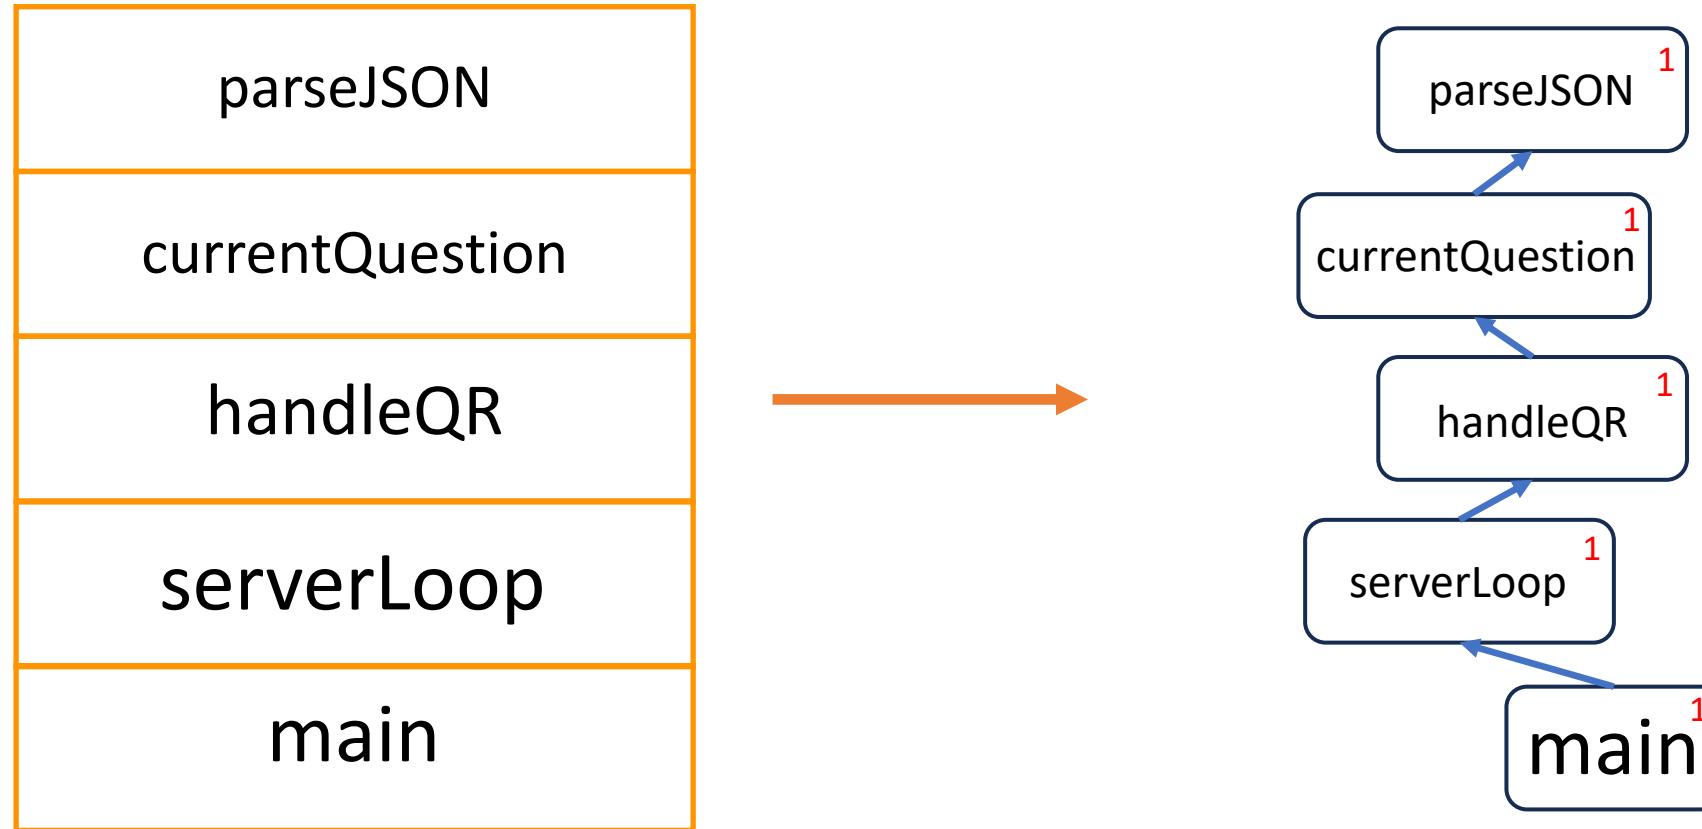


Turning traces into flamegraphs

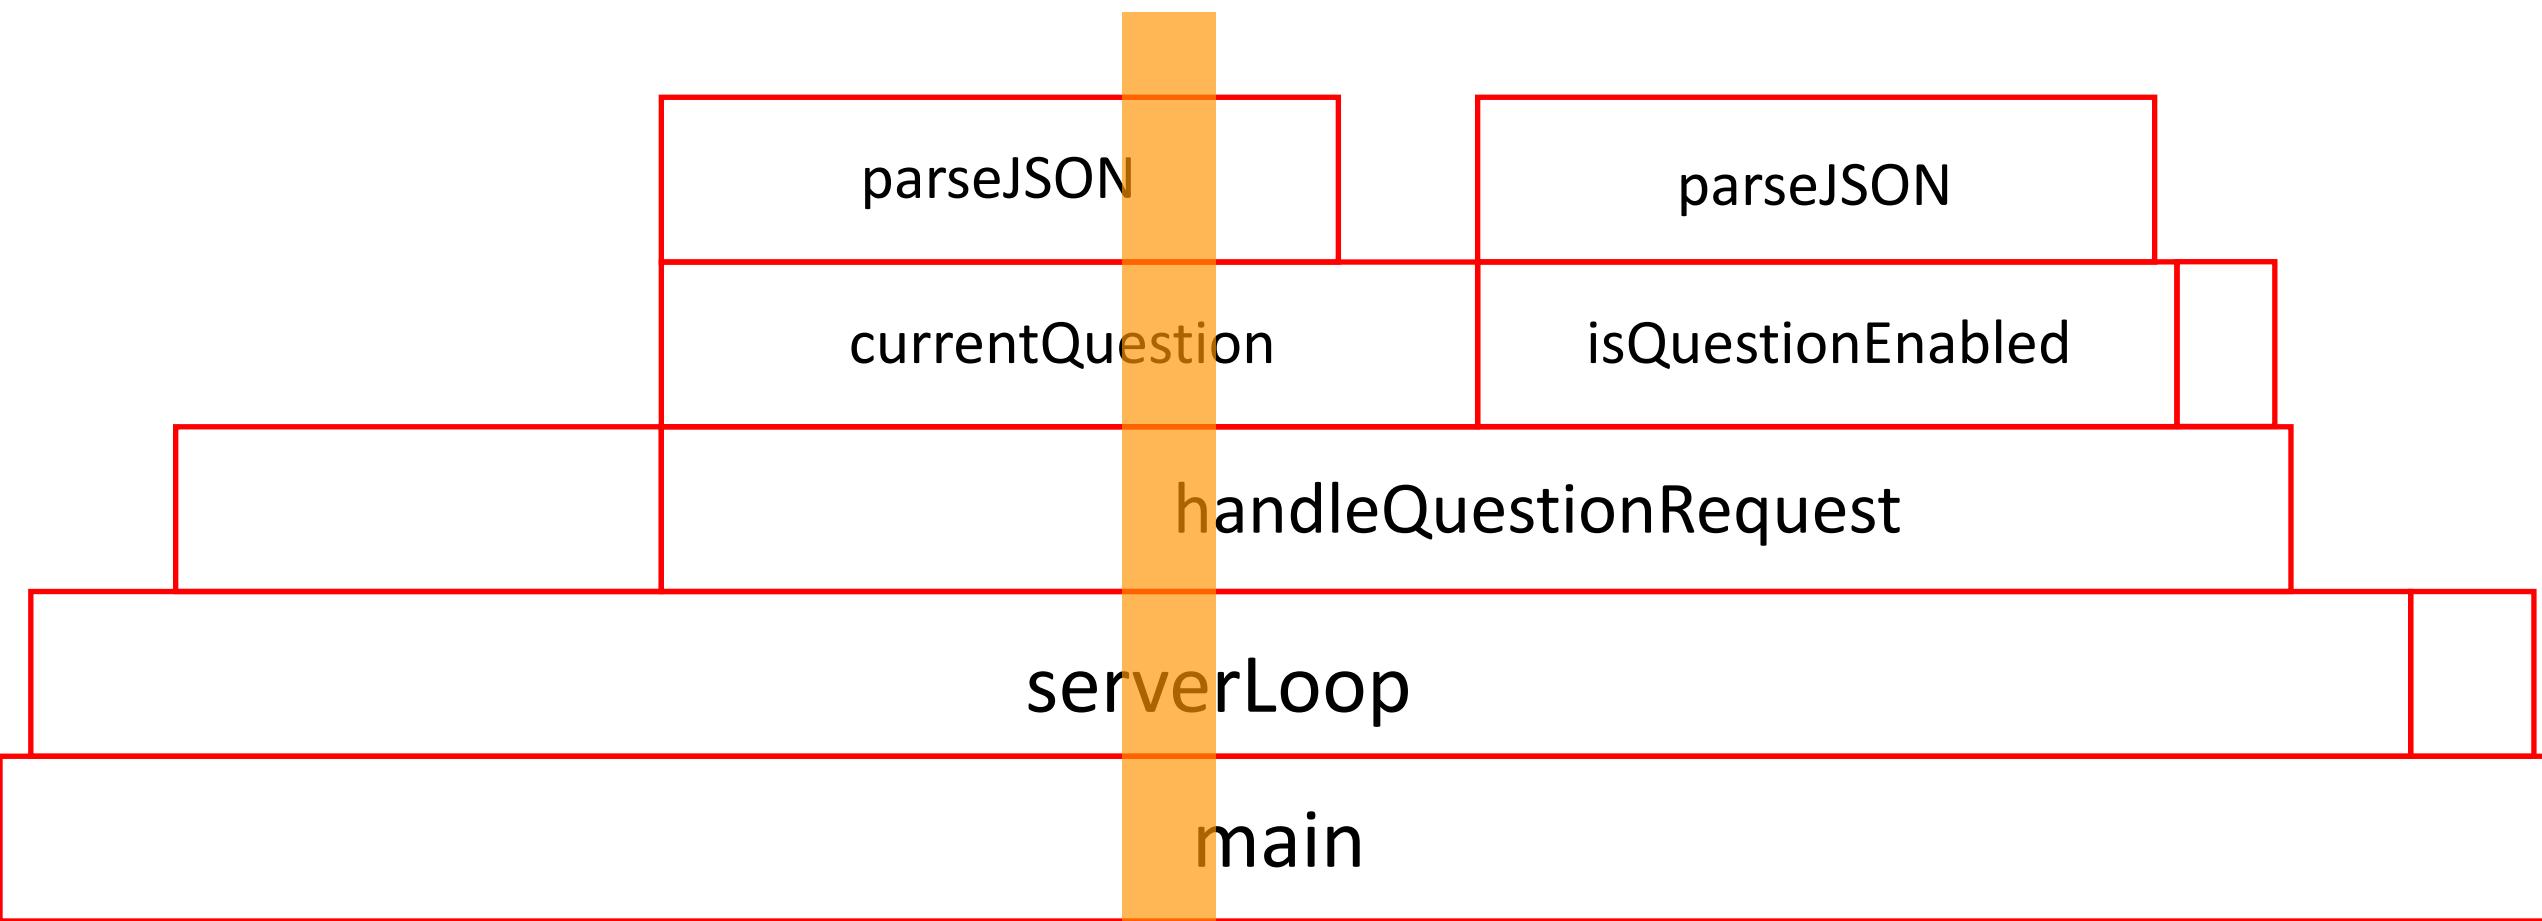


```
private static class Node {  
    final String method;  
    final Map<String, Node> children;  
    long samples = 0;  
  
    void addTrace(List<String> trace)  
}
```

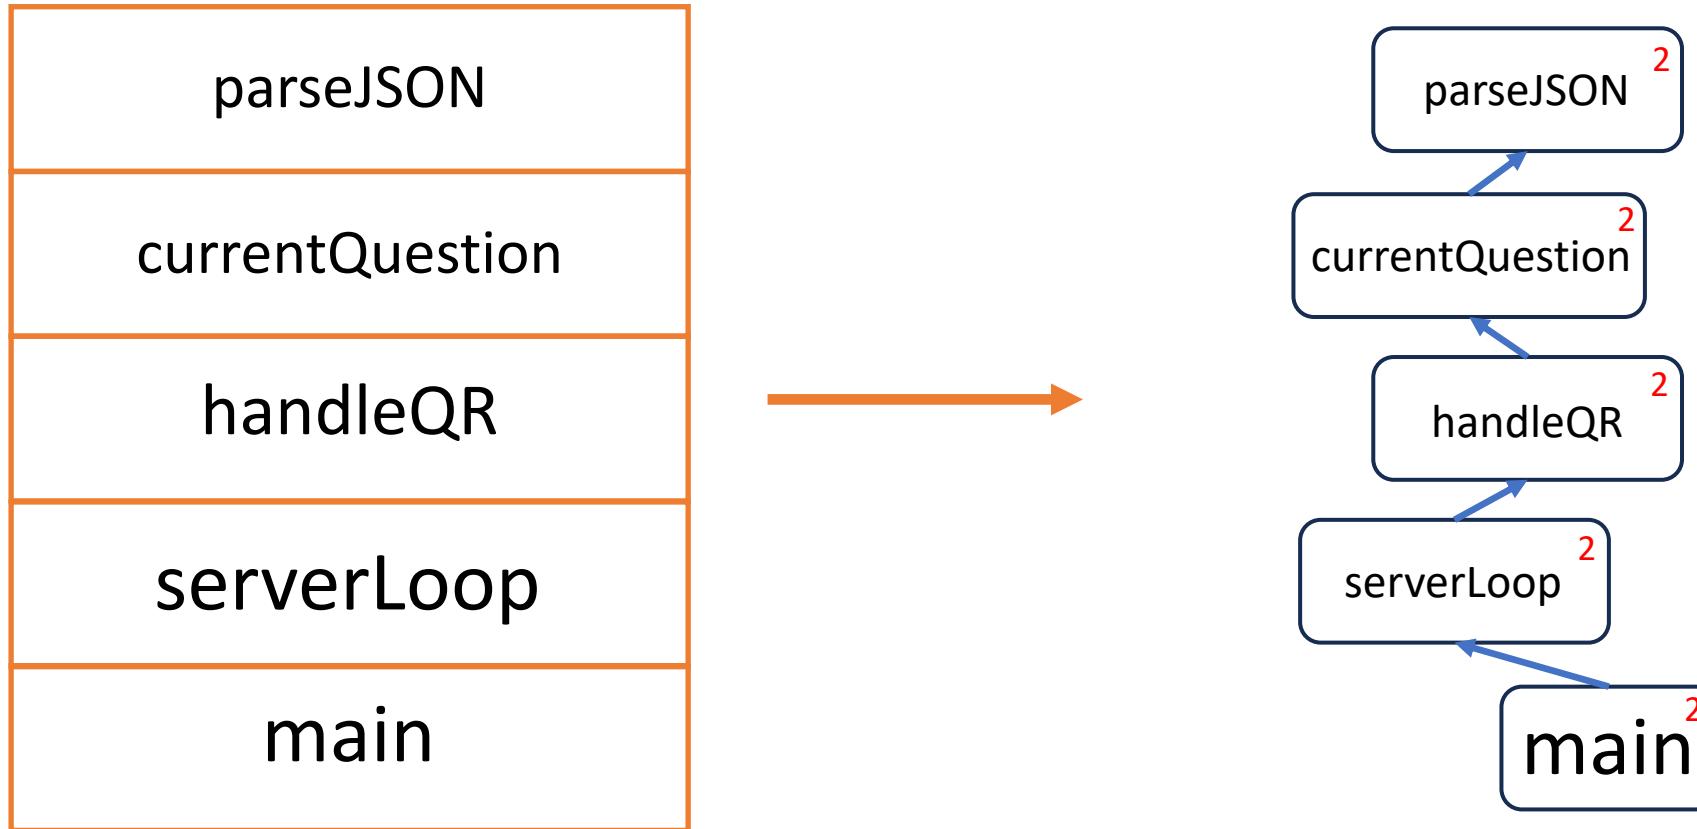
Turning traces into flamegraphs



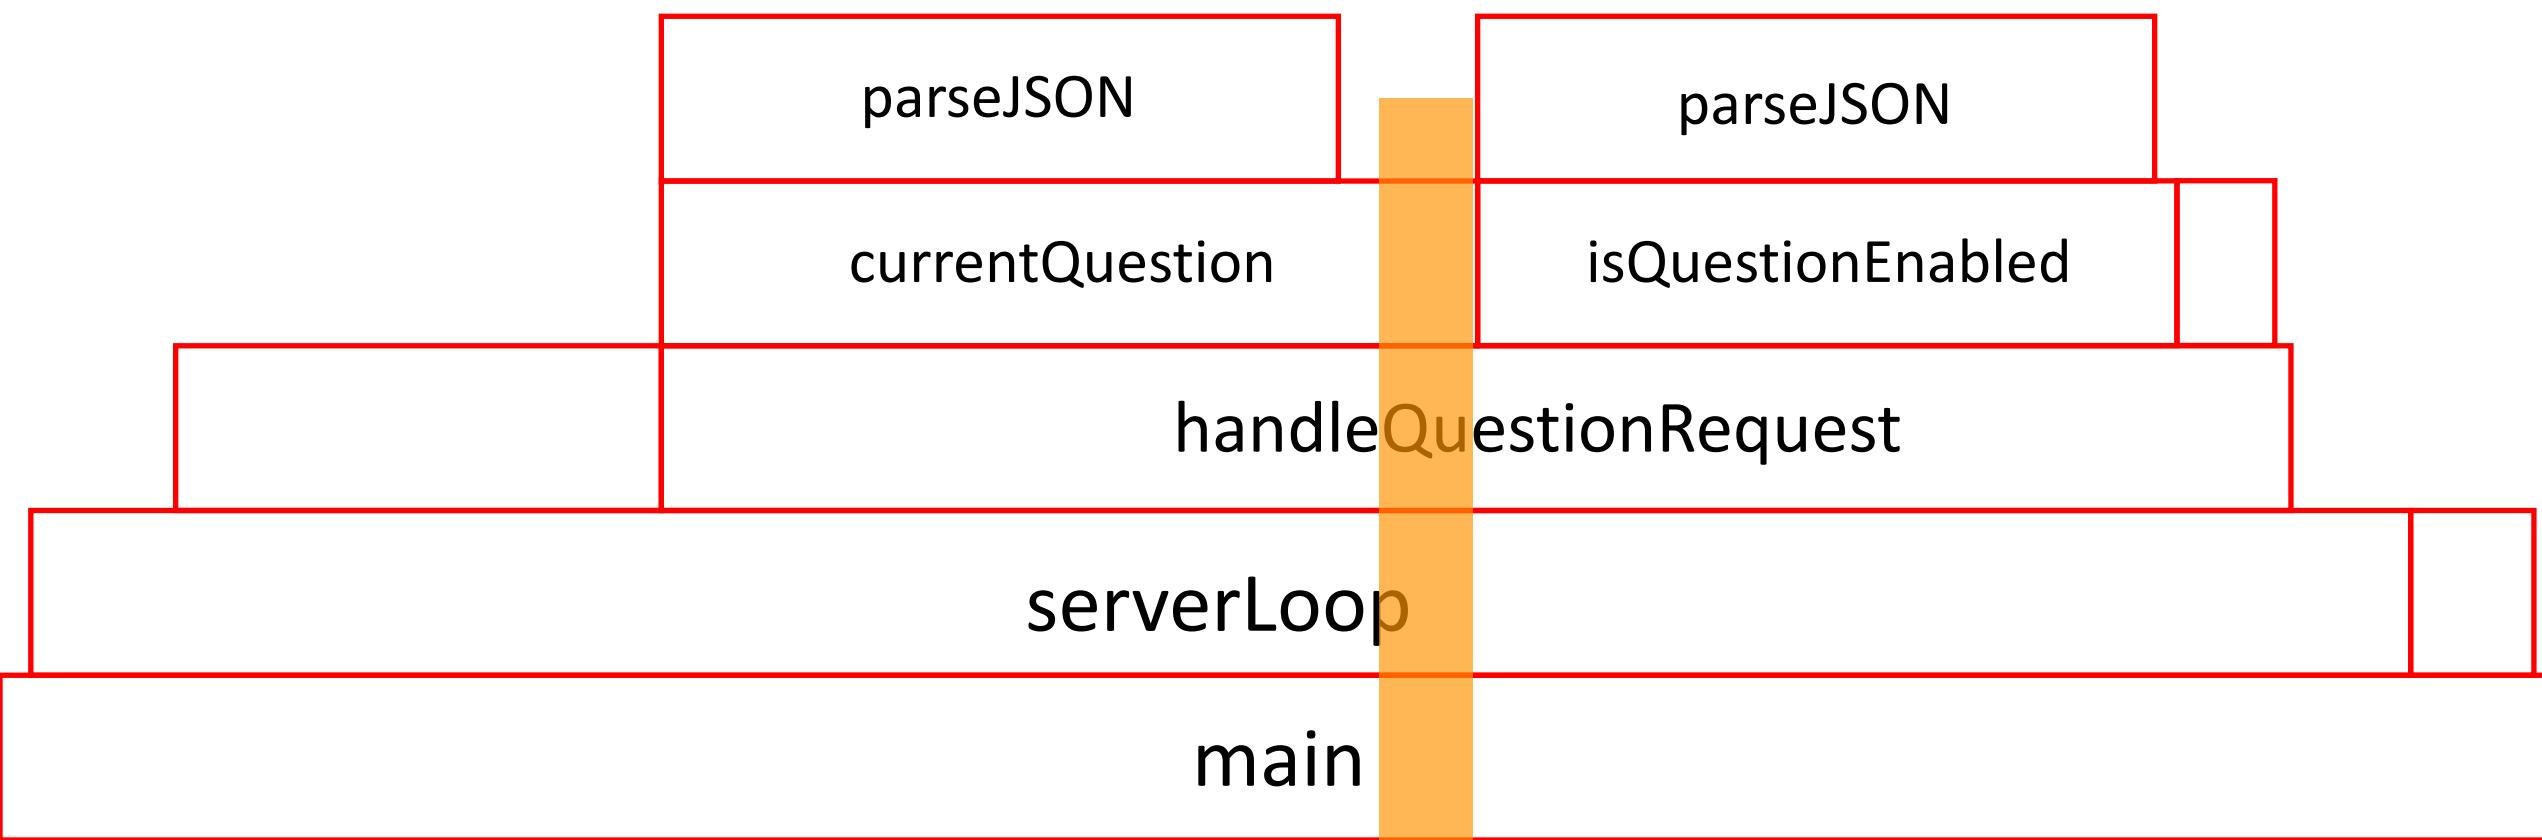
Flamegraphs would be great



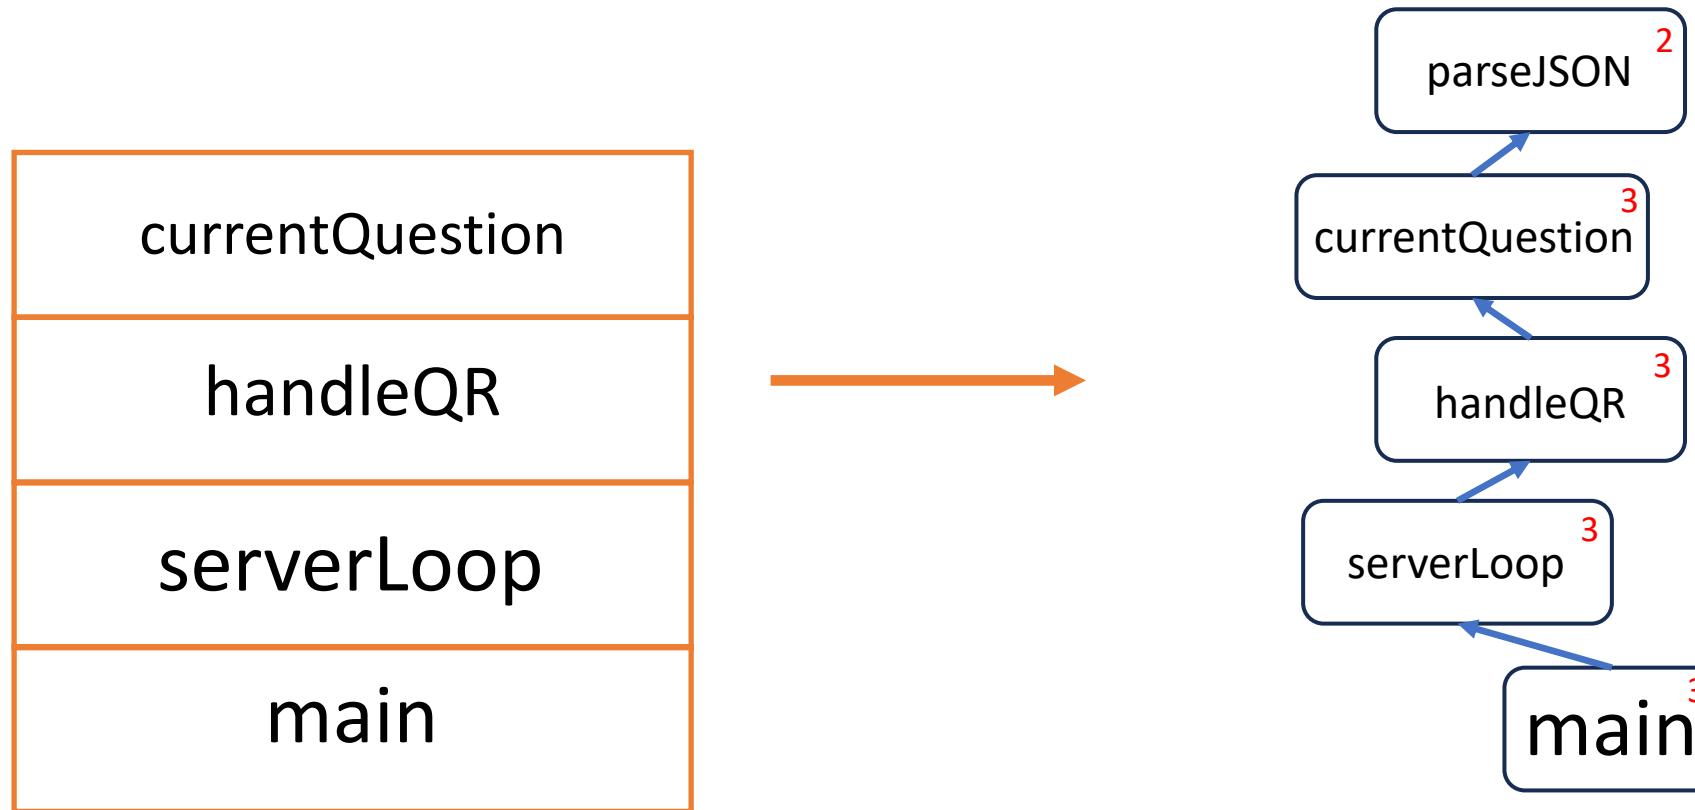
Turning traces into flamegraphs



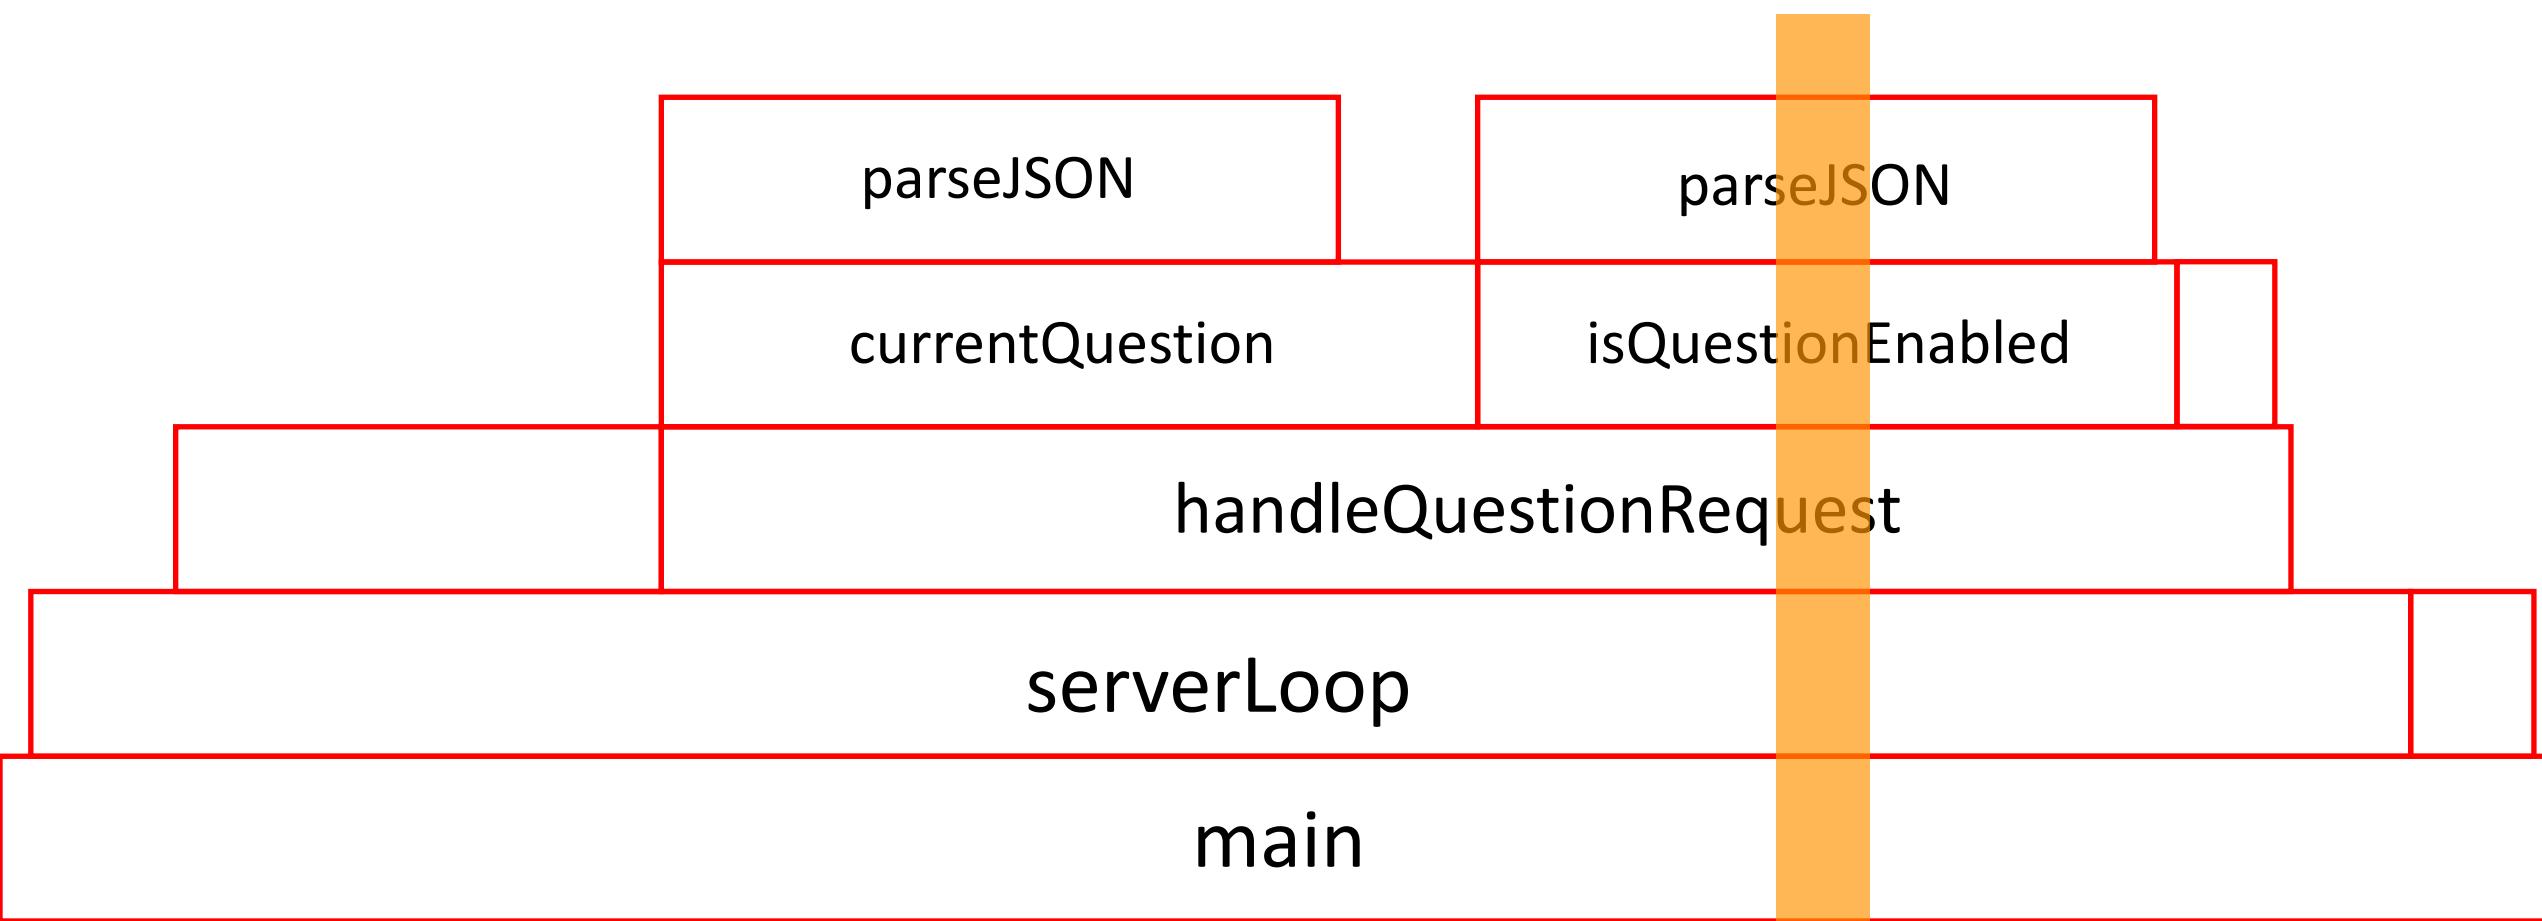
Flamegraphs would be great



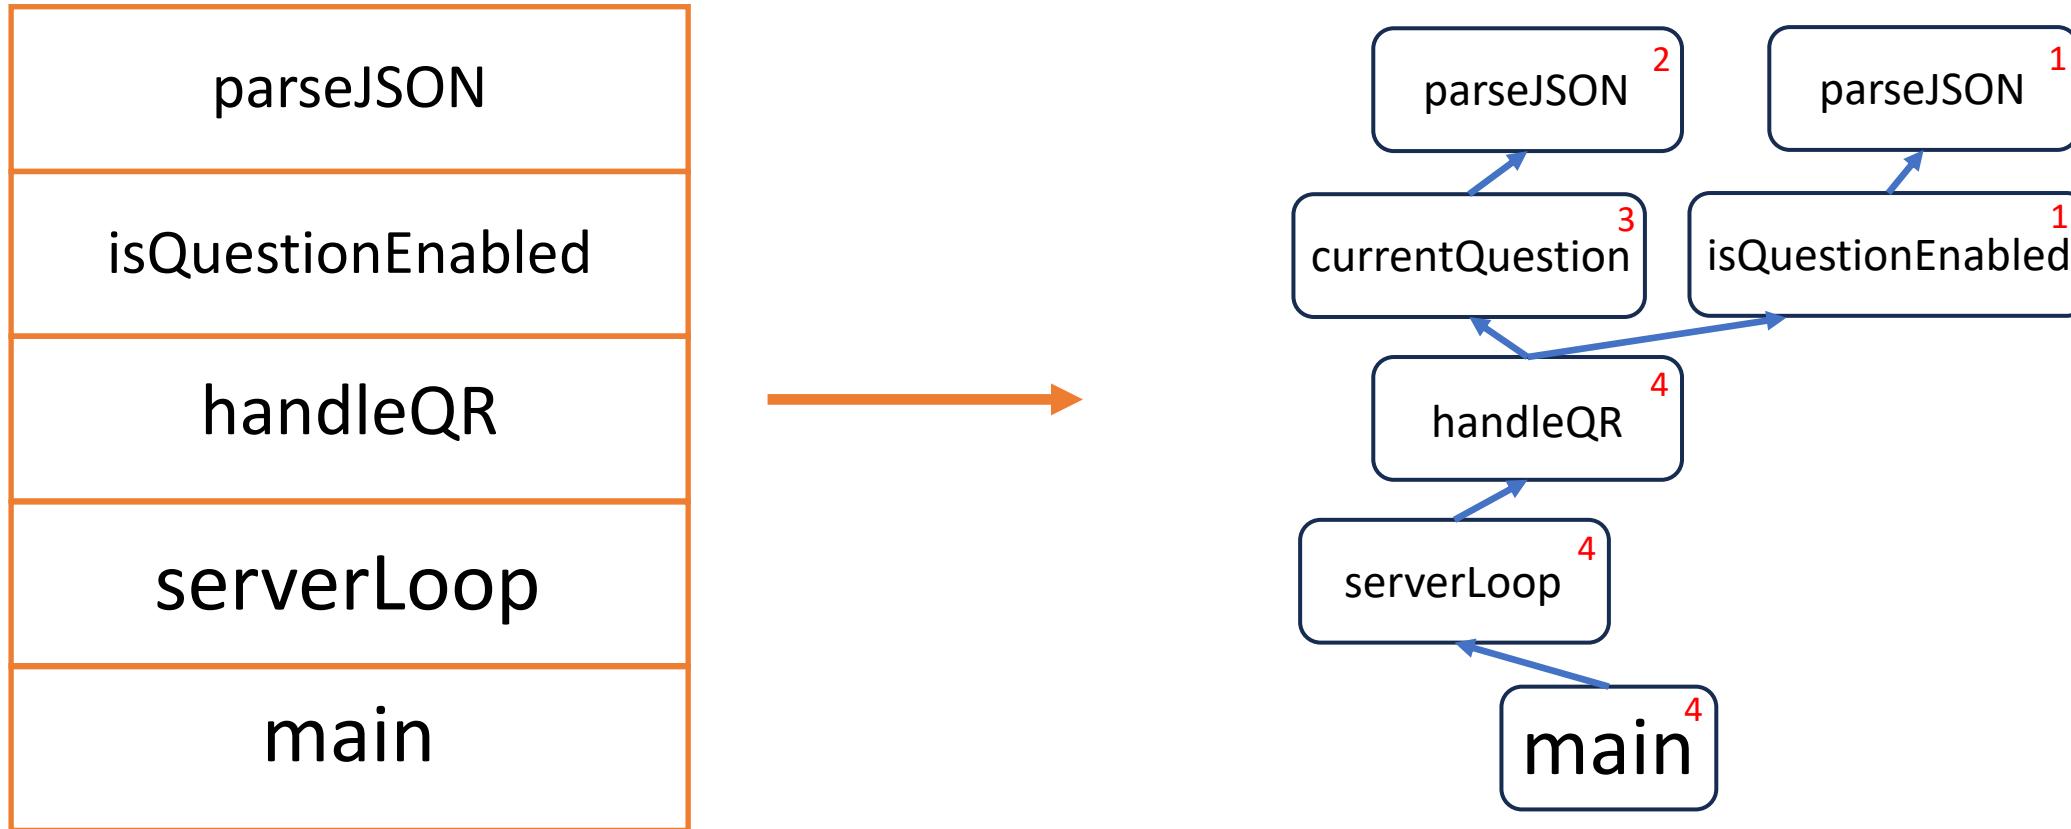
Turning traces into flamegraphs



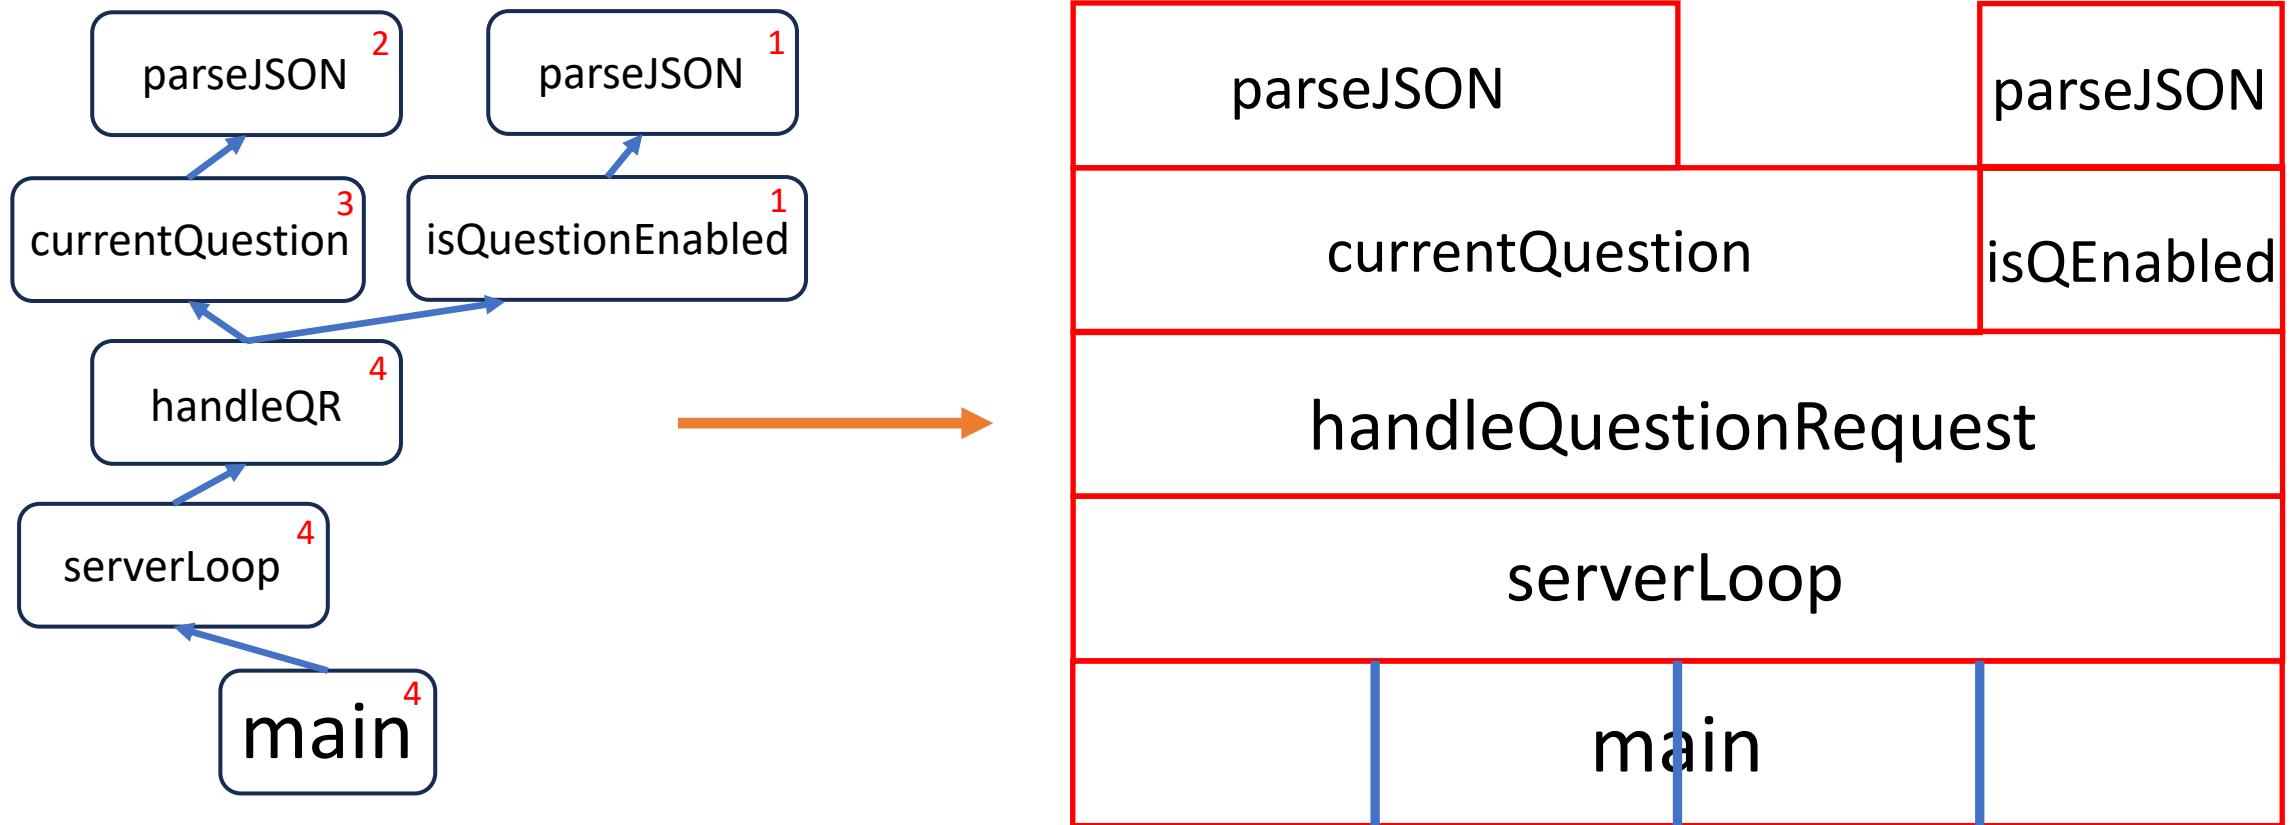
Flamegraphs would be great



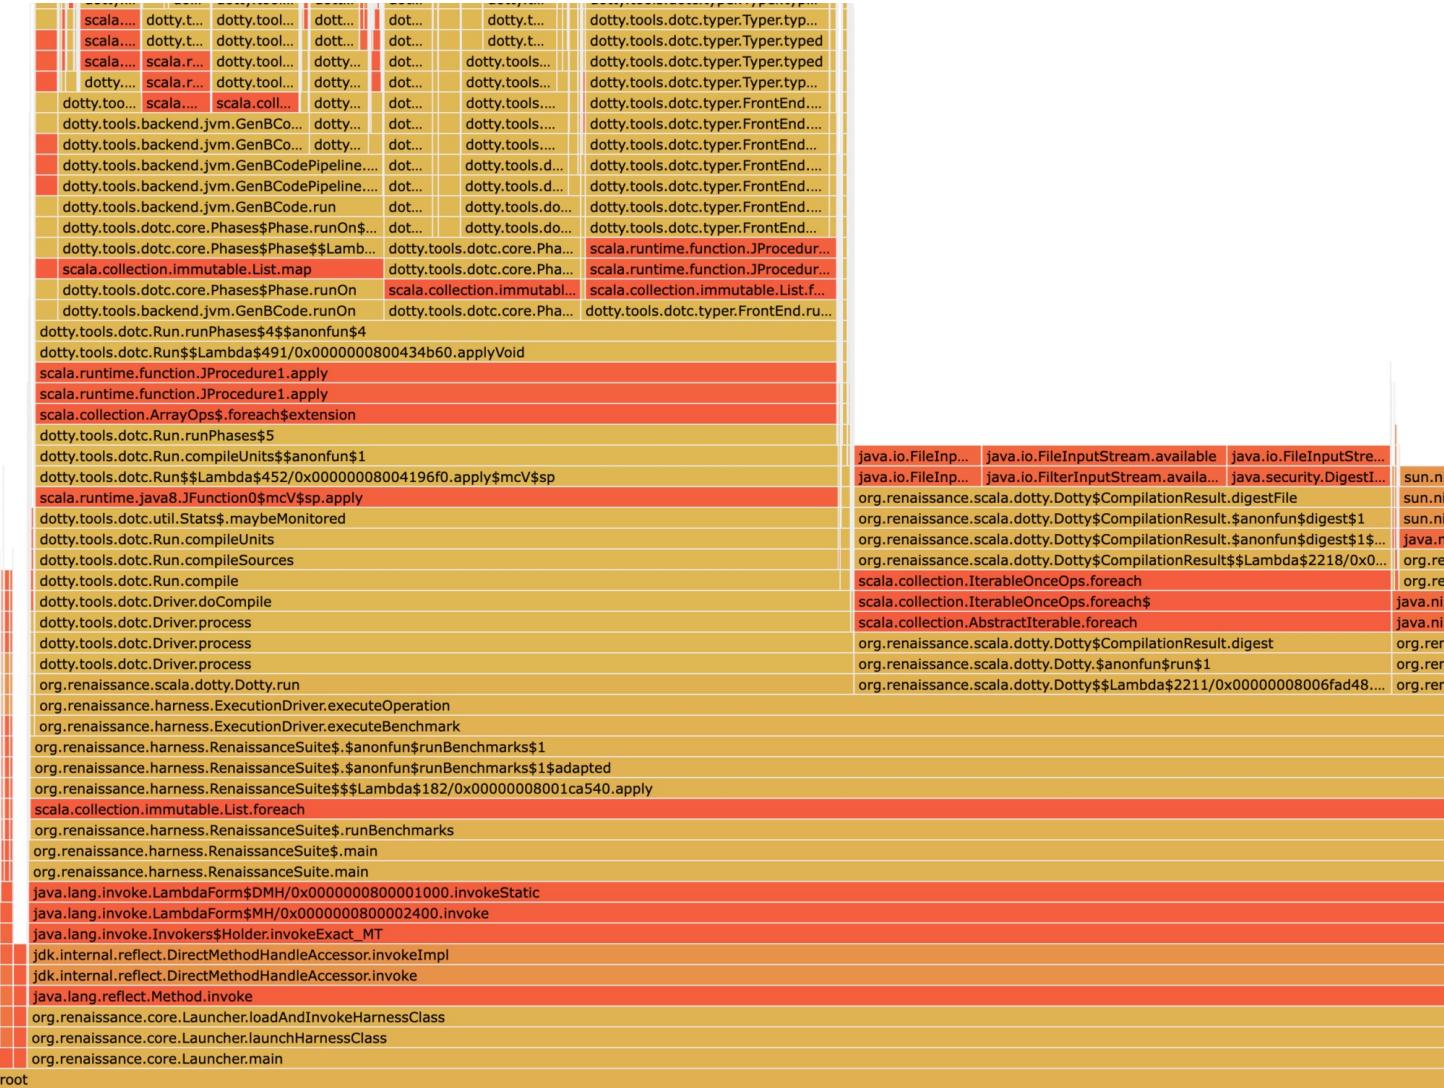
Turning traces into flamegraphs



Turning traces into flamegraphs

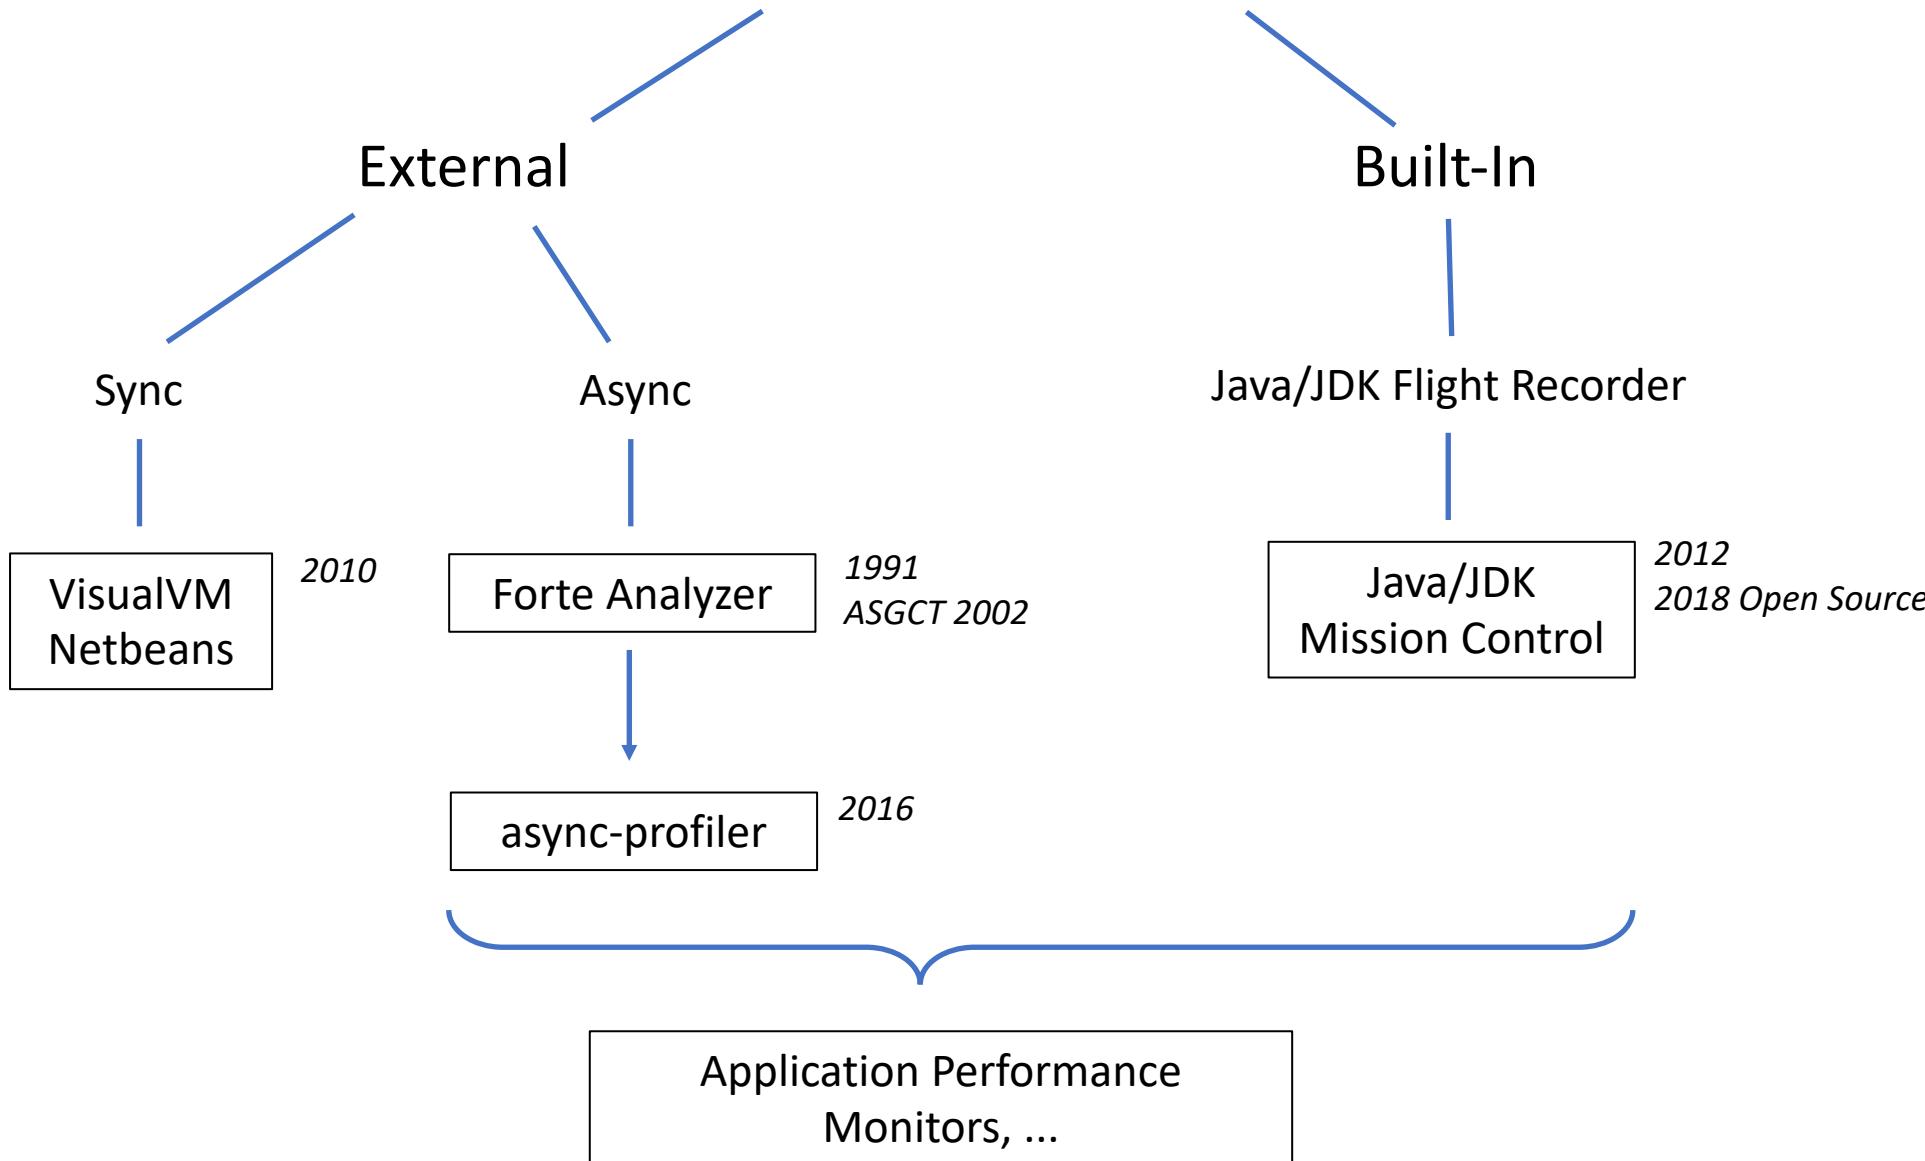


With the power of d3-flame-graph



Reality

Sampling Profiler



Obtaining a profile

Async Profiler

```
java \
  -agentpath:libasyncProfiler.so=start,\
  event=cpu,\
  file=flame.html,flamegraph \
arguments
```

Download from [GitHub](#)

JDK Flight Recorder (JFR)

```
java \
  -XX:+UnlockDiagnosticVMOptions \
  -XX:+DebugNonSafepoints \
  -XX:+FlightRecorder \
  -XX:StartFlightRecording=filename=file.jfr \
arguments
```

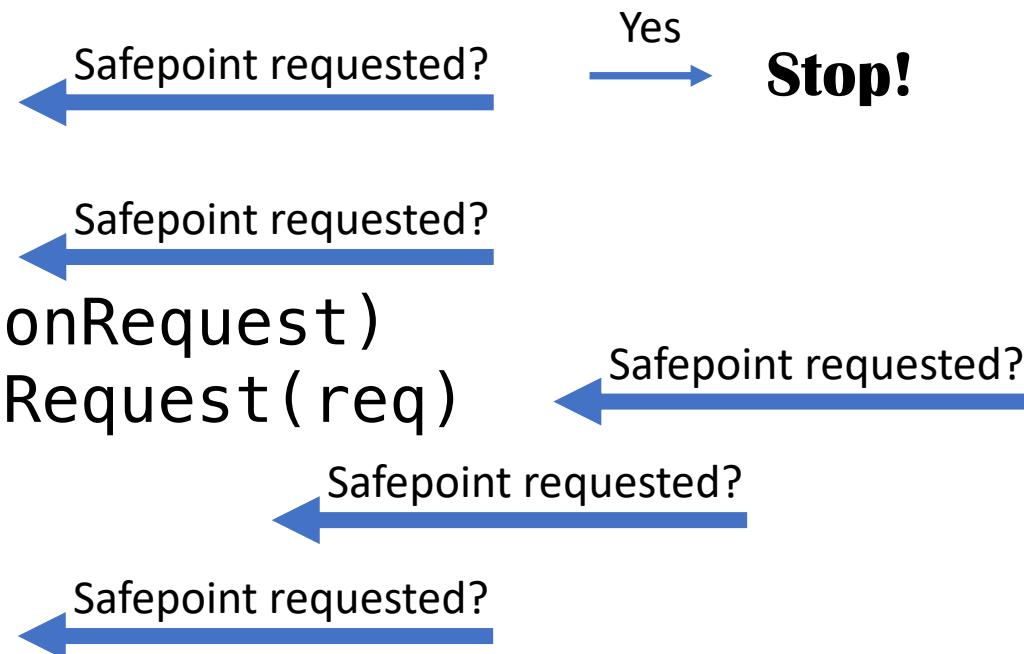
Already included in your JDK 8+

Don't trust them



Safepoint

```
fun serverLoop() {  
    while (true) {  
        req = ...  
        if (req.isQuestionRequest)  
            handleQuestionRequest(req)  
        ...  
    }  
}
```



Safepoint

Safepoint bias

Please stop,
please, when
you're ready.

Stop!

using Safepoints

fully asynchronous

Safepoint bias with multiple threads

Please stop,
please, when
you're ready.

Stop!
Thread 1

using Safepoints

fully asynchronous

AsyncGetCallTrace

ucontext →

Java frames

C/C++ frames

write
writeBytes
FileOutputStream::write
BufferedOutputStream::flushBuffer
BufferedOutputStream::implFlush
BufferedOutputStream::flush
PrintStream::implWrite
PrintStream::write
...
PrintStream::println
BasicSample::waitForEver
BasicSample::main
JavaCalls::call_helper
jni_invoke_static
...
_pthread_start

Stack

Stack walking



lineno (BCI)	25
method_id	&FileOutputStream::write

frames

num_frames
16

ASGCT_CallTrace

Tests in the OpenJDK?

1

But its a good one?

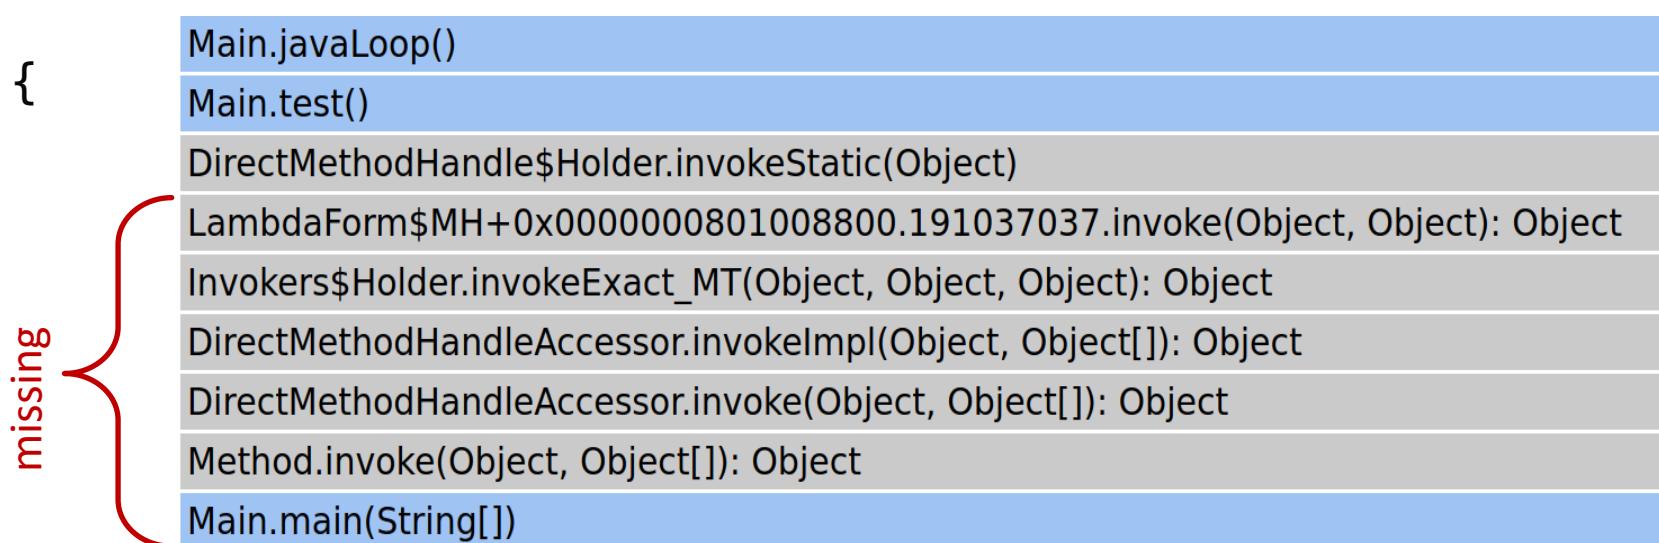
```
public class ASGCTBaseTest {  
    // ...  
    public static void main(String[] args) {  
        if (!check()) {  
            throw ...;  
        }  
    }  
    private static native boolean check(); ← checked just this method  
}
```

Take profiles with a
grain of salt

Profilers have bugs, e.g. JDK-8302320

```
public class Main {  
    public static void main(String[] args) ... {  
        Class<?> klass = Main.class;  
        Method mainMethod = klass.getMethod("test");  
        mainMethod.invoke(null);  
    }  
    static void test() {  
        javaLoop();  
    }  
    static void javaLoop() {  
        /* endless loop */  
    }  
}
```

Flamegraph



But surely OpenJ9 is
better?

0

Tests for AsyncGetCallTrace

Or GetStackTrace?

2

Safepoint biased is bad

```
@Benchmark
public void blameSetResult() {
    byte b = 0;
    for (int i = 0; i < size; i++) {
        b += buffer[i];
    }
    setResult(b); /* LINE 38 */
}

private void setResult(byte b) {
    setResult(b == 1); /* LINE 90 */
}

@CompilerControl(CompilerControl.Mode.DONT_INLINE)
private void setResult(boolean b) {
    result = b;
}
```

Profilers are just
software

Tests could better

> 1

Performance and accuracy
could be better

Could be safer*

*crashes are rare in practice

Not designed with safety in mind

8284828: Use `os::ThreadCrashProtection` to protect AsyncGetCallTrace from crashing #8225

closed

parttimenerd wants to merge 9 commits into `openjdk:master` from `parttimenerd:parttimenerd_8284828`

“ I should add that the CrashProtection mechanism was mainly put in place as a result of having to deliver JFR from JRockit into Hotspot under a deadline, upholding feature-parity. The stack walking code was in really bad shape back then.

– Markus Grönlund

“ AGCT is a legacy mechanism that was created for one single unsupported purpose

– David Holmes

TDD

Test

Driven

Development

TD

TD

R D D

Rage

Driven

Development

More Tests

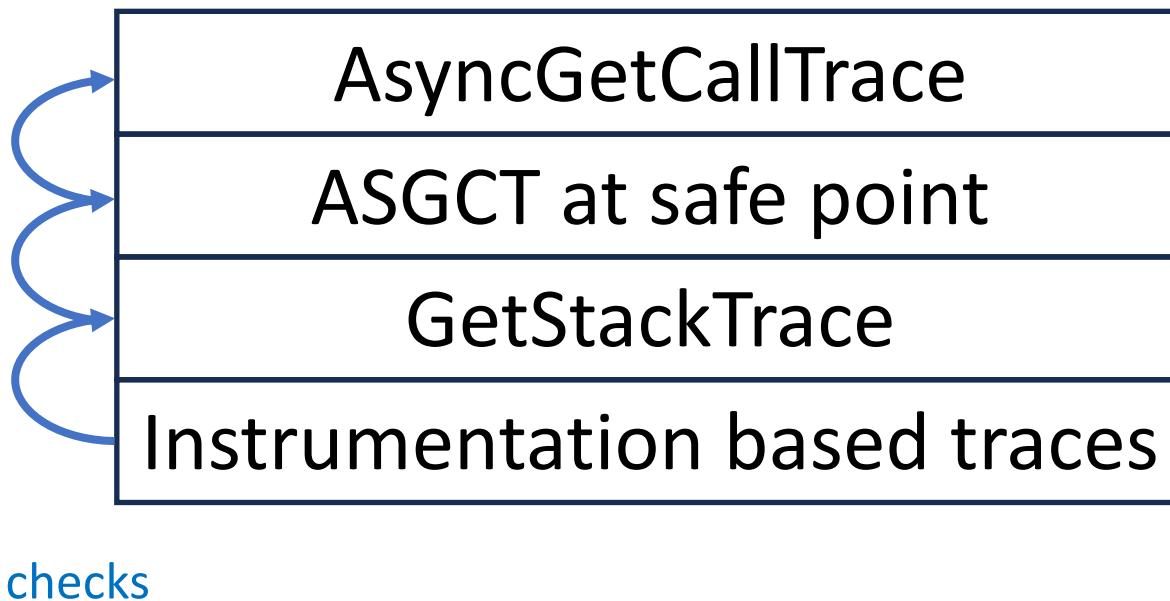






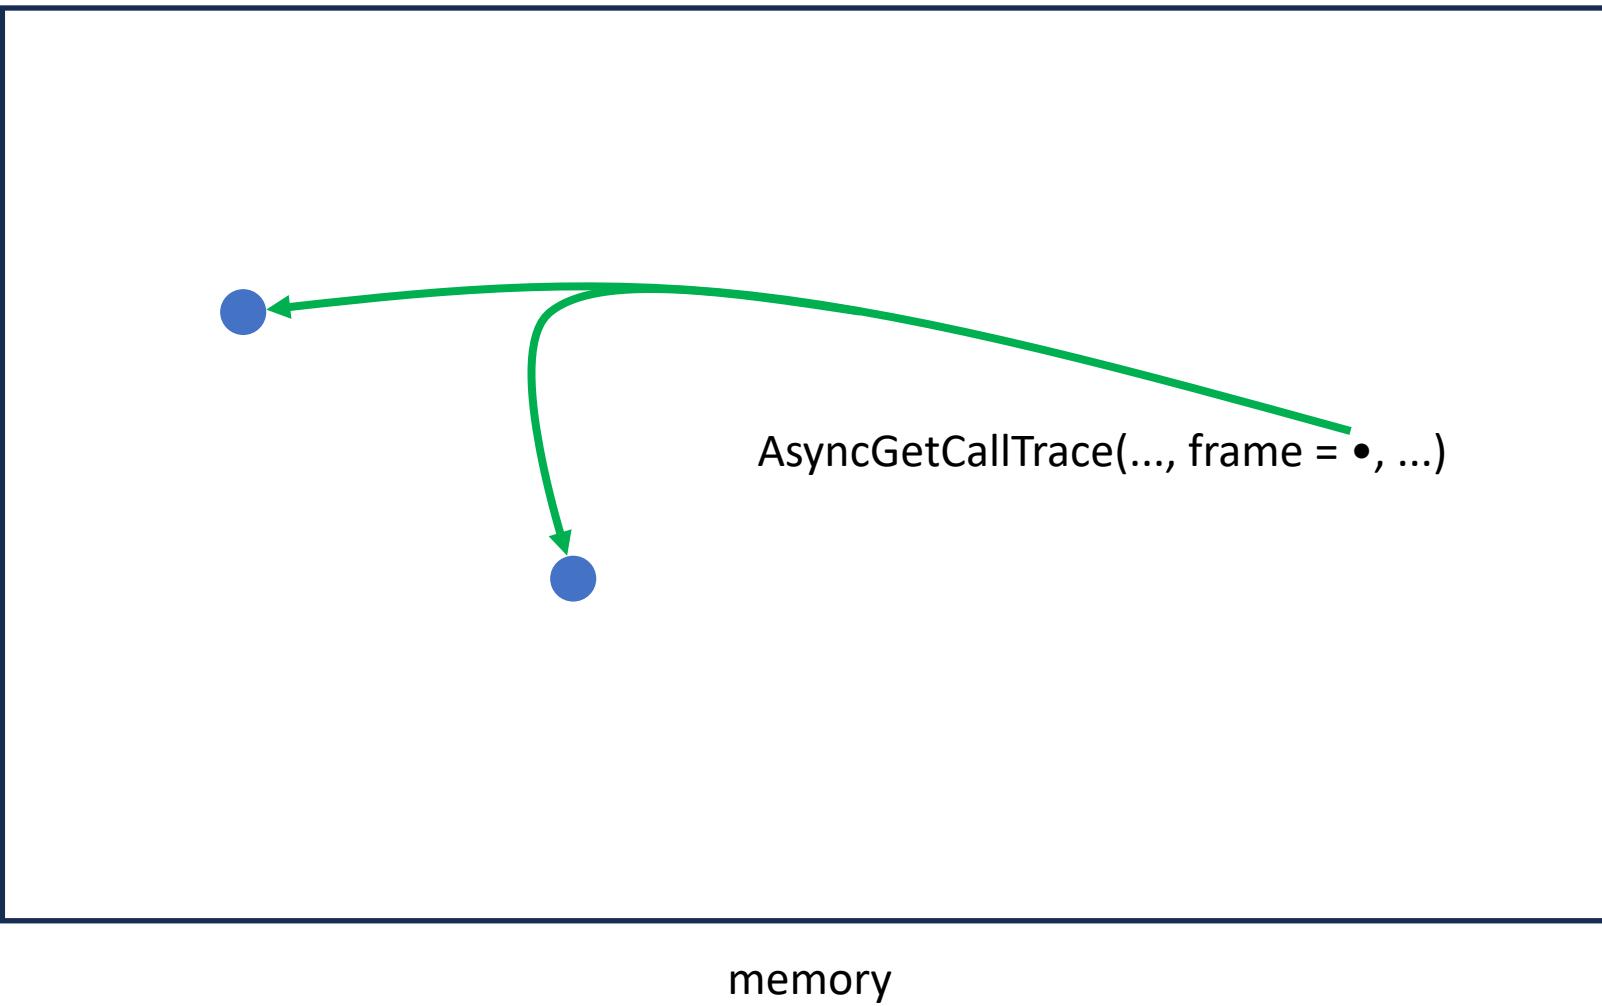
Testing Approaches

Layered Oracle



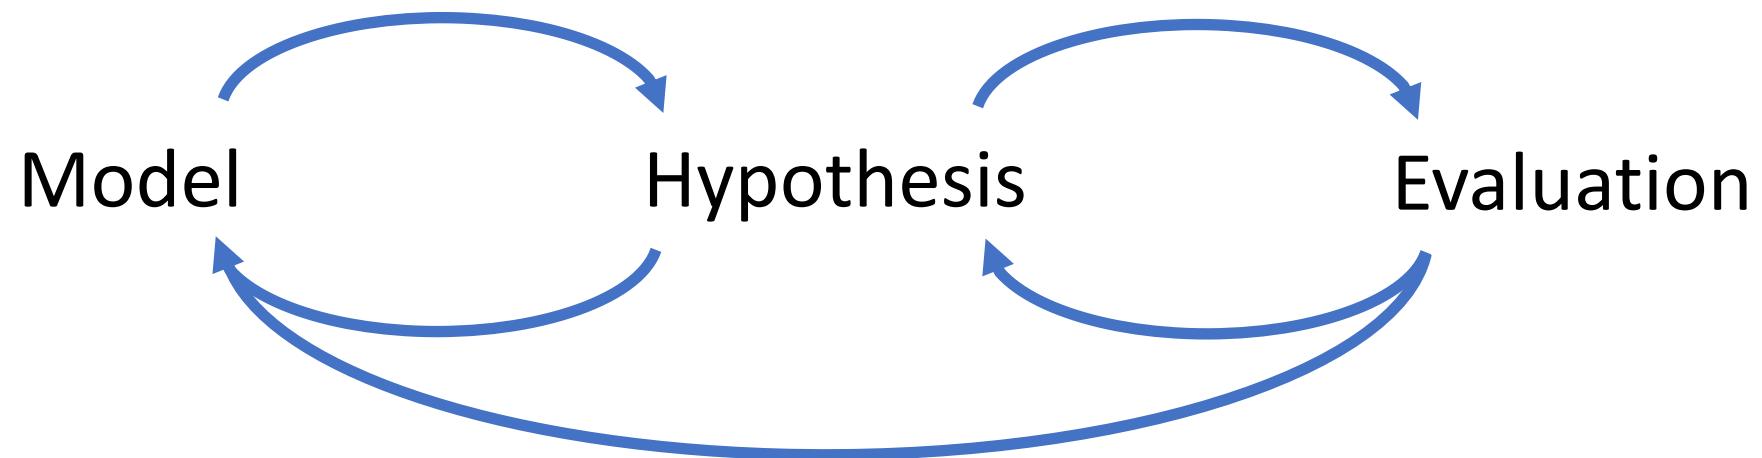
```
fun serverLoop() {  
    logEntry("serverLoop")  
    while (true) {  
        req = ...  
        if (req.isQuestionRequest)  
            handleQuestionRequest(req)  
        ...  
    }  
    logExit("serverLoop")  
}
```

Fuzzing for safety



Experimentation technique

Profiling Loop



Profilers are great

But don't trust them too much

@mostlynerdless.de on BlueSky
parttimenerd on GitHub
mostlynerdless.de

@SweetSapMachine on Twitter
sapmachine.io

